

CS660: Intro to Database Systems

Class 10: Log-Structured-Merge Trees

Instructor: Manos Athanassoulis

<https://bu-disc.github.io/CS660/>

Reads vs Writes: The two extremes


Assume **no index** – what is the **best way to physical store** our data?



Case 1: I have a static datasets and I **only receive reads**

how to read?

binary search




Sorted!



Case 2: I **only receive new updates**, which I never try to read

scan




Append (log)


How to bridge the two?

Consider a workload with **bursts of new data**, followed by queries!

Append:



Once we accumulate “enough” data, we sort, and we write to the disk



What to do if we still receive incoming data?

Keep the sorted file



&


append to a new one








What to do with many sorted files?


Merge them!
































LSM-tree

The Log-Structured Merge-Tree (LSM-Tree)

1996

Patrick O'Neil¹, Edward Cheng²
Dieter Gawlick³, Elizabeth O'Neil¹
To be published: Acta Informatica

Patrick O'Neil
UMass Boston




LSM-tree
O'Neil *et al.*

1996

✓ good sequential reads & writes

✓ good random writes



array
of discs

why?

RAID, striping ← ?

✗ LSM not explicitly needed

LSM-tree
O'Neil *et al.*

so, arrays of disks were enough!



how many IOPS?

10KRPM

max seek time 1.5ms

100 disks

10KRPM: 10K rev in 60s

$60/10000 = 6\text{ms}$ per rev

avg. rot. delay: 3ms (6ms/2)

avg. seek time: 0.75ms (1.5ms/2)

1 I/O / 3.75ms: 267 IOPS

100 disks: 26,700 IOPS



Bigtable

1980s


1996

2006

a decade

✓ good sequential reads & writes


✓ good random writes



array
of discs

✗ worse sequential access

✗ bad random writes



commodity
hardware

what happened in 2006?



LSM-tree
O'Neil *et al.*

We set up a Bigtable cluster with N tablet servers to measure the performance and scalability of Bigtable as N is varied. The tablet servers were configured to use 1 GB of memory and to write to a GFS cell consisting of 1786 machines with two 400 GB IDE hard drives each.



Bigtable

1980s


1996

2006

a decade

✓ good sequential reads & writes

✓ good random writes



array
of discs


SSD wear-friendly

competitive rand. reads

fast ingestion (sequential)

✗ worse sequential access

✗ bad random writes



commodity
hardware



LSM-tree
O'Neil *et al.*

1980s

1996

2006

a decade



Bigtable

LSM-tree
O'Neil *et al.*

1996



Bigtable

2006

APACHE
HBASE 

2007

LSM-tree
O'Neil *et al.*

1996



2006



2007



2010

LSM-tree
O'Neil *et al.*

1996



2006



2007



2010



2011

LSM-tree
O'Neil *et al.*

1996




Bigtable

2006



2007



cassandra

2010



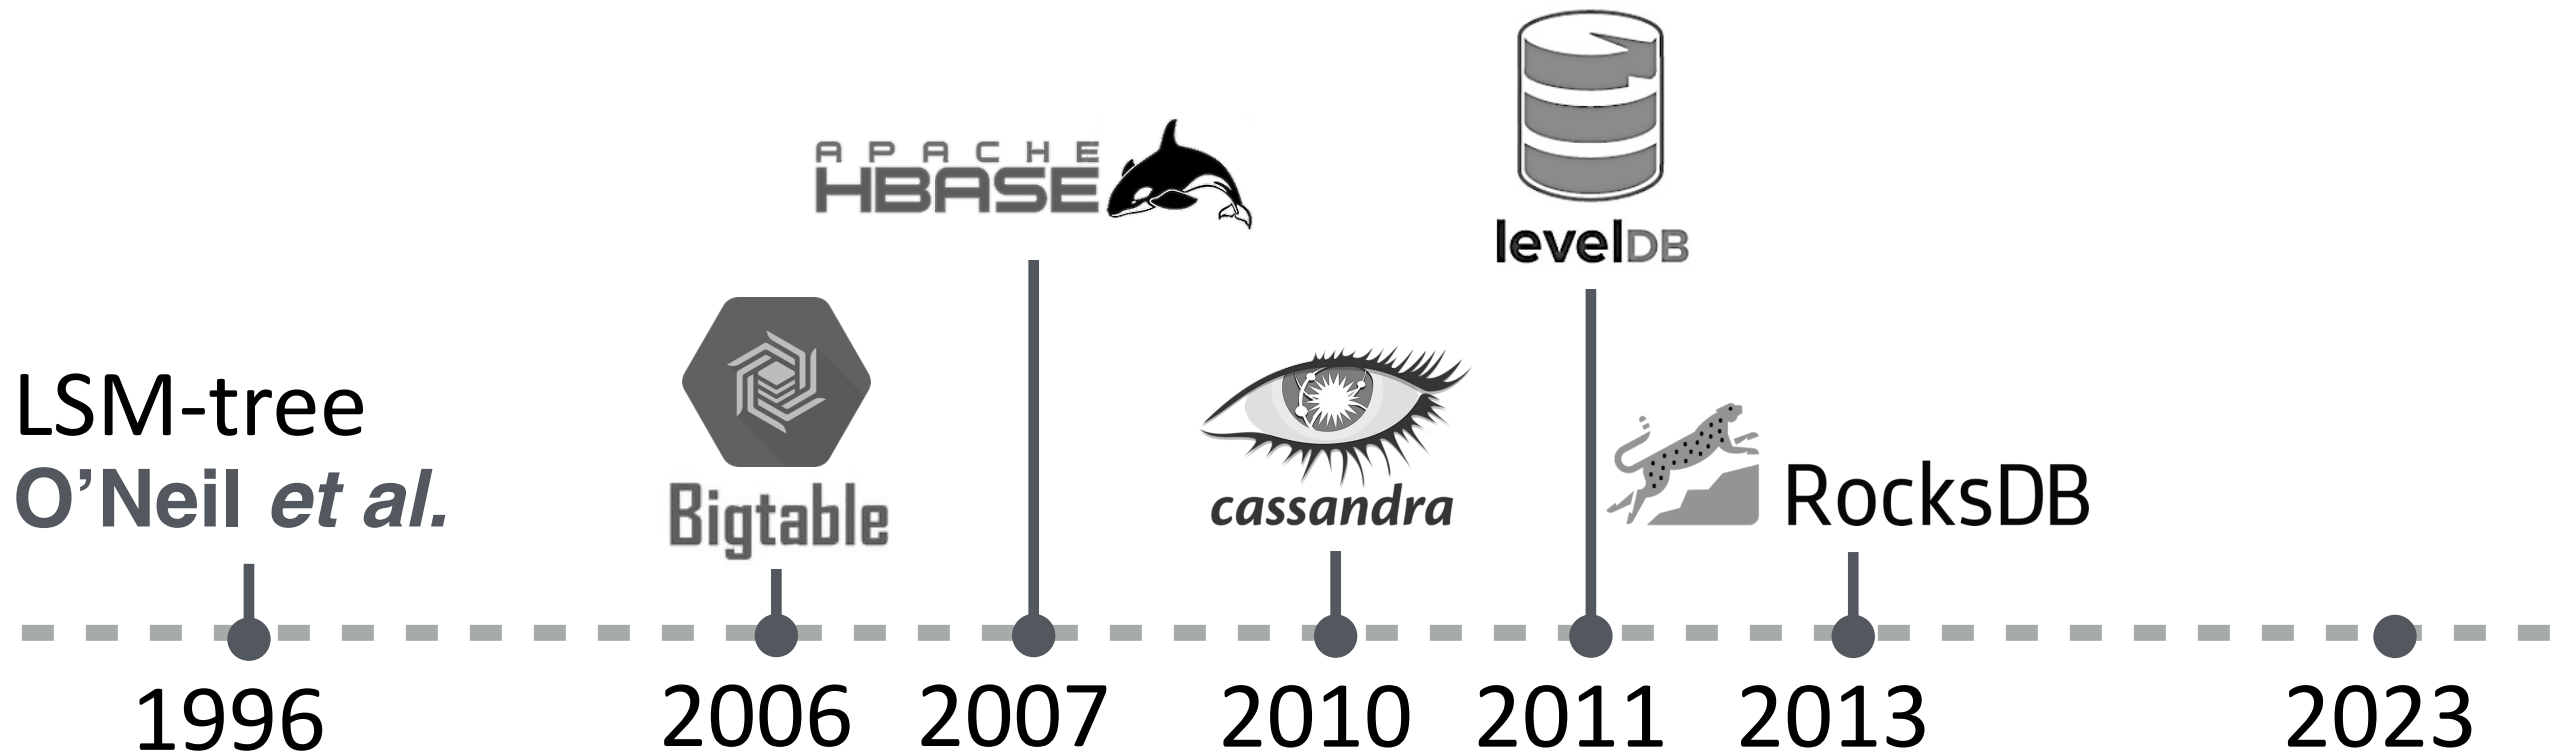
levelDB

2011



RocksDB

2013



LSM-tree

NoSQL



RocksDB




levelDB



SCYLLA



cassandra



tarantool




Bigtable

APACHE
HBASE



DynamoDB



SQLite



relational



influxdb



QuasarDB

time-series

2023

LSM-tree

NoSQL



RocksDB




levelDB



SCYLLA



cassandra



tarantool




Bigtable

APACHE
HBASE



DynamoDB



SQLite



relational



influxdb



QuasarDB

time-series

2023

How does LSM-tree compare with prior approaches?

Compare and contrast data structures.

What to use when?

Data Structure	Lookup cost	Insertion cost
Sorted array		
Log		
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue

Compare and contrast data structures.

What to use when?

Data Structure	Lookup cost	Insertion cost
Sorted array		
Log		
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Sorted Array

Measure Performance in I/Os

n entries

B entries fit into a disk block


Array spans $N = \frac{n}{B}$ disk blocks

Lookup method & cost?

Binary search: $O(\log_2(N))$ I/Os

Insertion cost?

Push entries: $O(N/2)$ I/Os



Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log		
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log		
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Log (append-only array)

n entries

B entries fit into a disk block

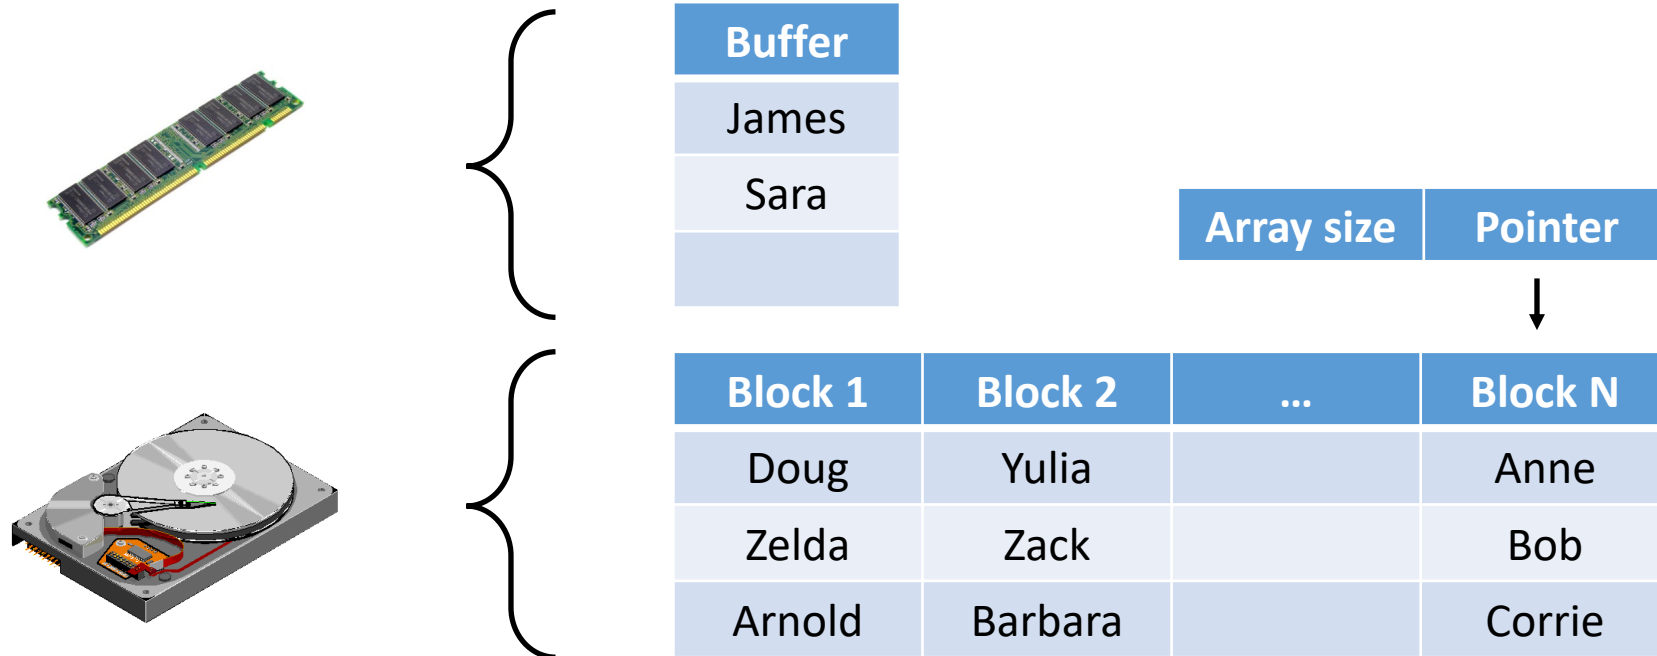
Array spans $N = \frac{n}{B}$ disk blocks

Lookup method & cost?

Scan: $O(N)$

Insertion cost?

Append: $O\left(\frac{1}{B}\right)$



Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree		
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

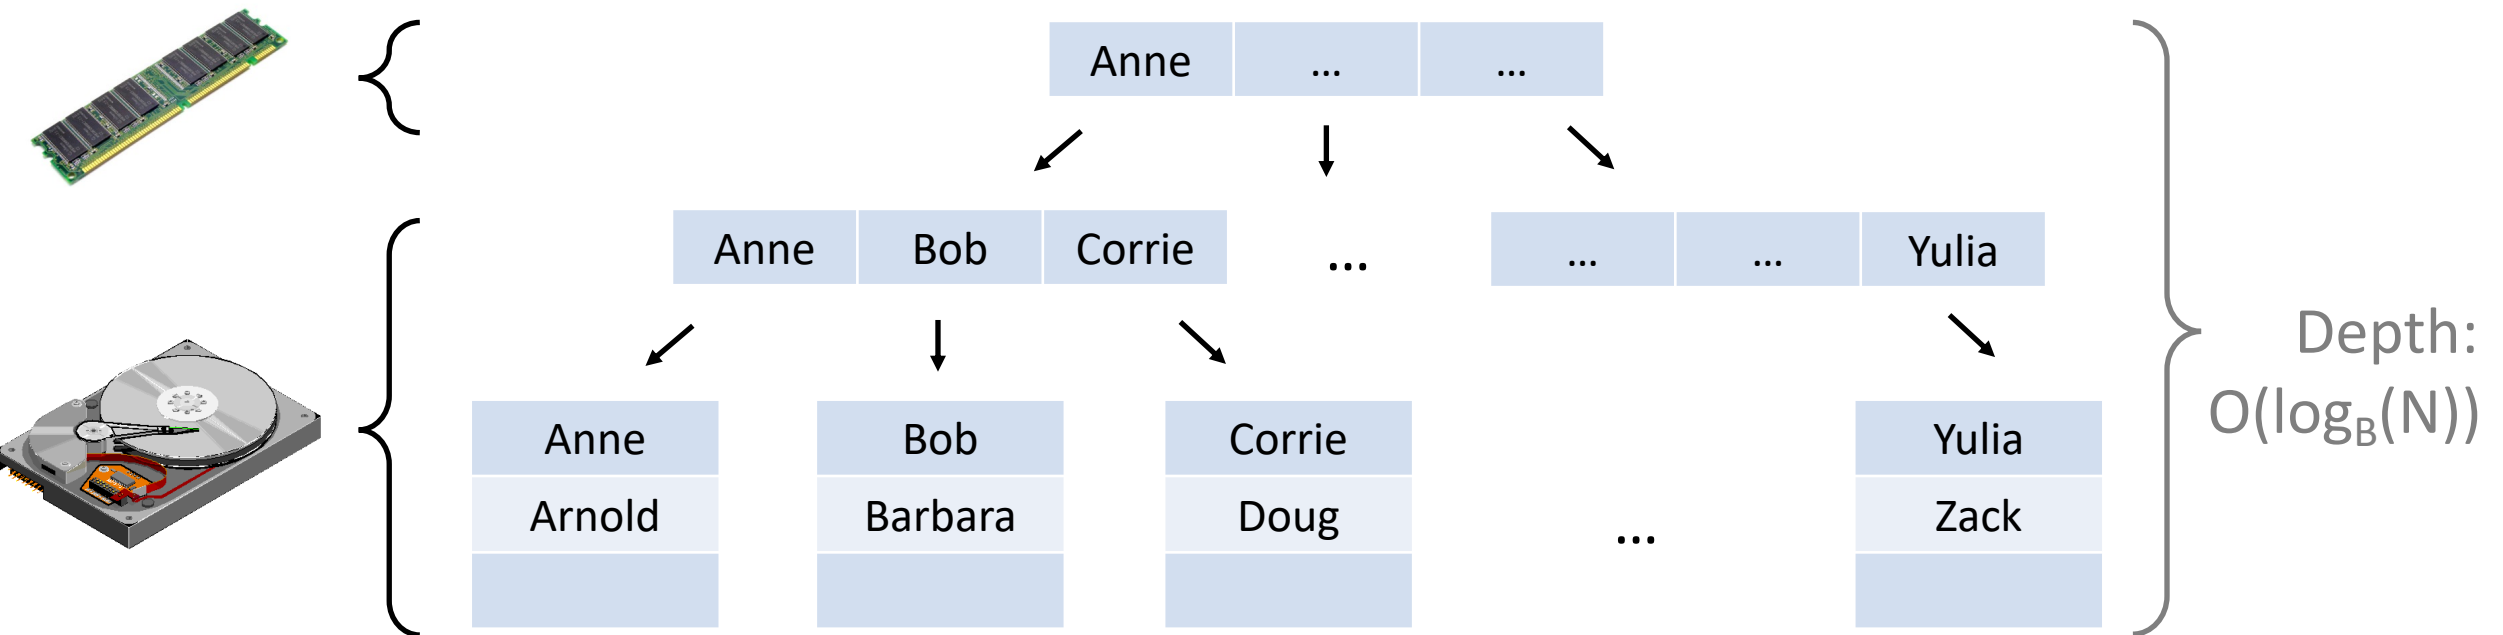
B-tree

Lookup method & cost?

Tree search: $O(\log_B(N))$

Insertion method & cost?

Tree search & append: $O(\log_B(N))$



Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

B-trees



Goetz Graefe

Microsoft, HP Fellow, now Google
ACM Software System Award


“It could be said that the world’s information is at our fingertips because of B-trees”

B-trees are no longer sufficient

Cheaper storage

Workloads more **insert-intensive**

We need **better insert-performance**



Results Catalogue


Goal to combine

sub-constant insertion cost
logarithmic lookup cost

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree		
Leveled LSM-tree		
Tiered LSM-tree		

Basic LSM-trees


Basic LSM-tree



Basic LSM-tree

Design principle #1:


optimize for insertions by buffering



Basic LSM-tree

Design principle #1:

optimize for insertions by buffering




Inserts



Basic LSM-tree

Design principle #1:


optimize for insertions by buffering



Basic LSM-tree

Design principle #1:

optimize for insertions by buffering




Basic LSM-tree

Design principle #1:


optimize for insertions by buffering

Design principle #2:


optimize for lookups by sort-merging arrays



Inserts



sort & flush buffer





Basic LSM-tree

Design principle #1:

optimize for insertions by buffering

Design principle #2:

optimize for lookups by sort-merging arrays



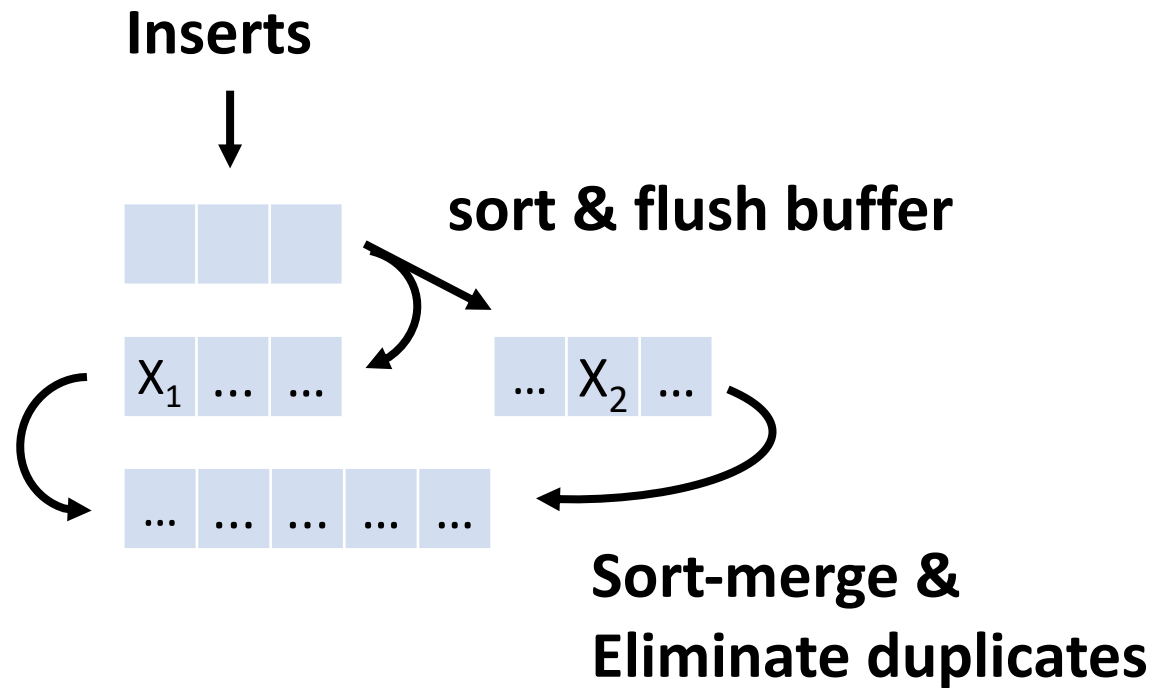

Basic LSM-tree

Design principle #1:

optimize for insertions by buffering

Design principle #2:

optimize for lookups by sort-merging arrays




Basic LSM-tree

Design principle #1:

optimize for insertions by buffering

Design principle #2:

optimize for lookups by sort-merging arrays




Basic LSM-tree

Design principle #1:


optimize for insertions by buffering

Design principle #2:


optimize for lookups by sort-merging arrays




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example



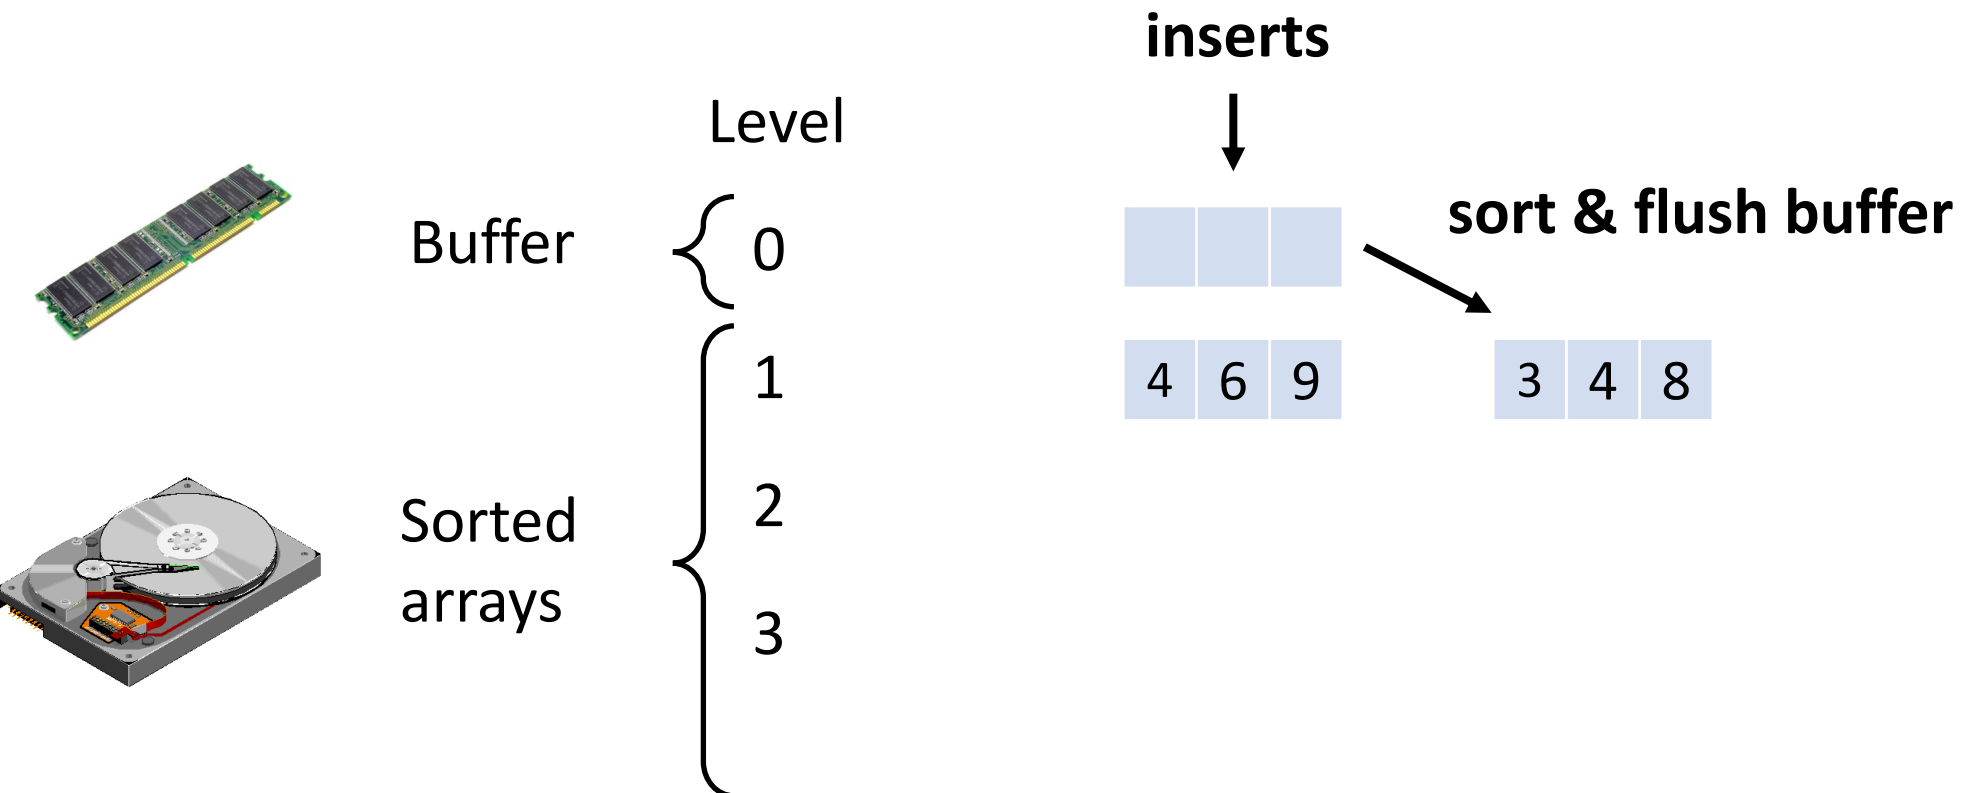
Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example



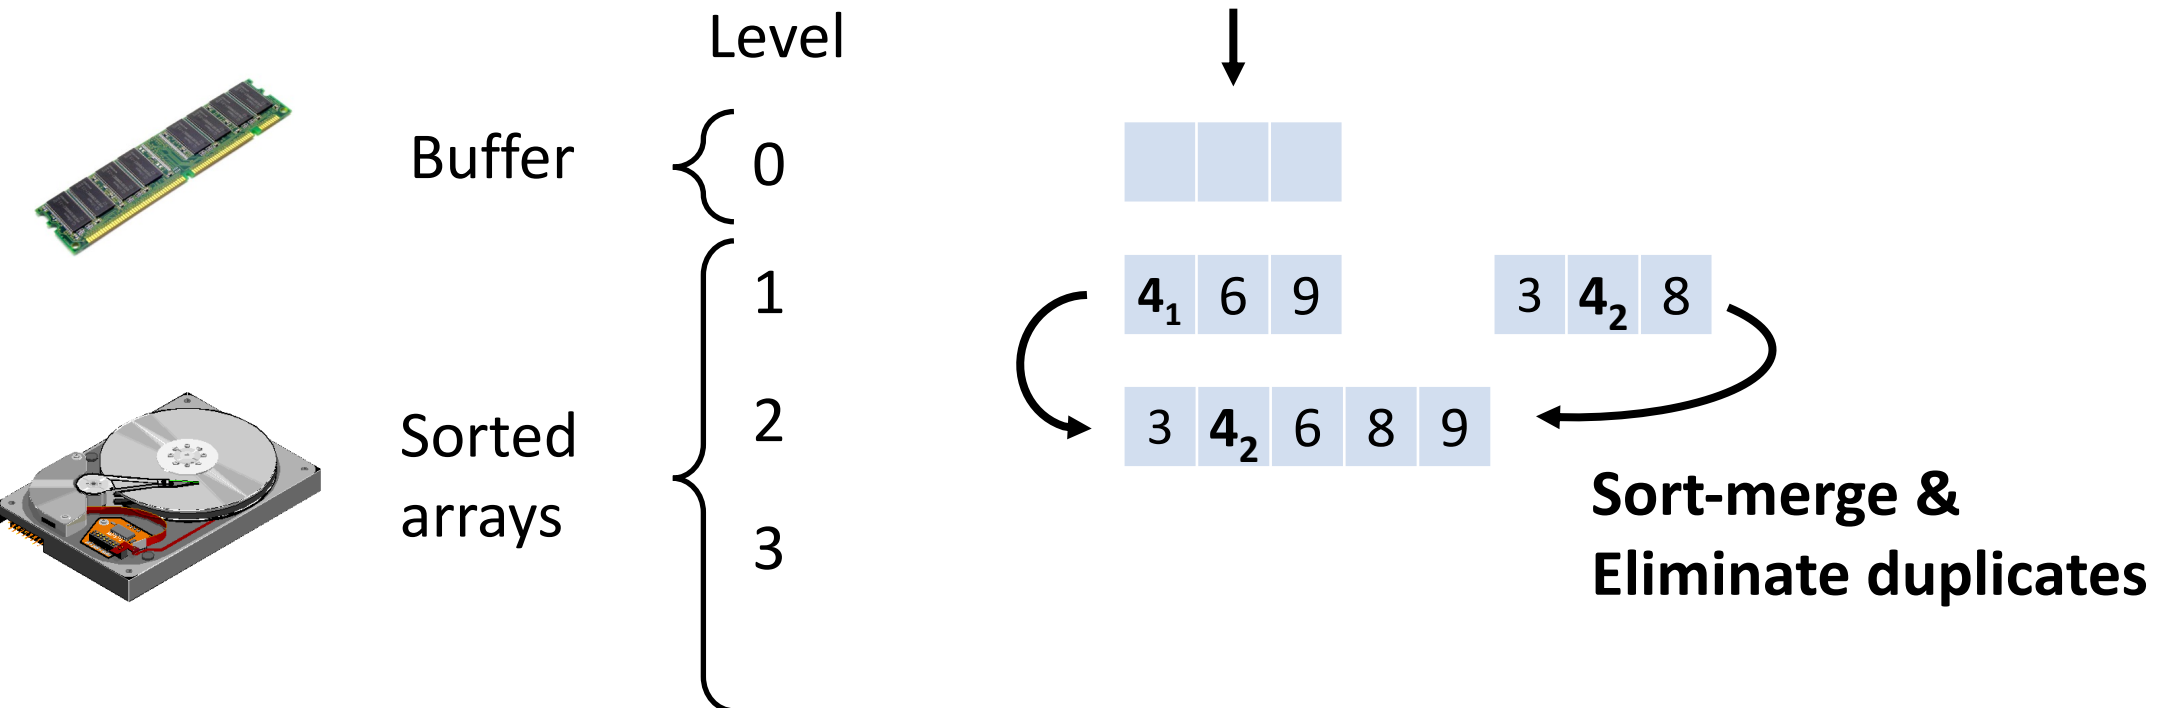
Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example




Basic LSM-tree – Example



Basic LSM-tree – Example



Basic LSM-tree – Example




Basic LSM-tree

Levels have exponentially increasing capacities.

How many levels?



$\log_2(N)$



Basic LSM-tree – Lookup cost

Lookup method?

How?

Lookup cost?


Search youngest to oldest.

Binary search.

$O(\log_2(N))$

$O(\log_2(N))$

$O(\log_2(N)^2)$



Basic LSM-tree – Insertion cost

How many times is each entry copied?


$O(\log_2(N))$, once per level

What is the price of each copy?

$O(1/B)$, amortized

Total insert cost?

$O((1/B) \cdot \log_2(N))$



Results Catalogue

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue

Better insert cost and **worse lookup cost** compared with B-trees

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

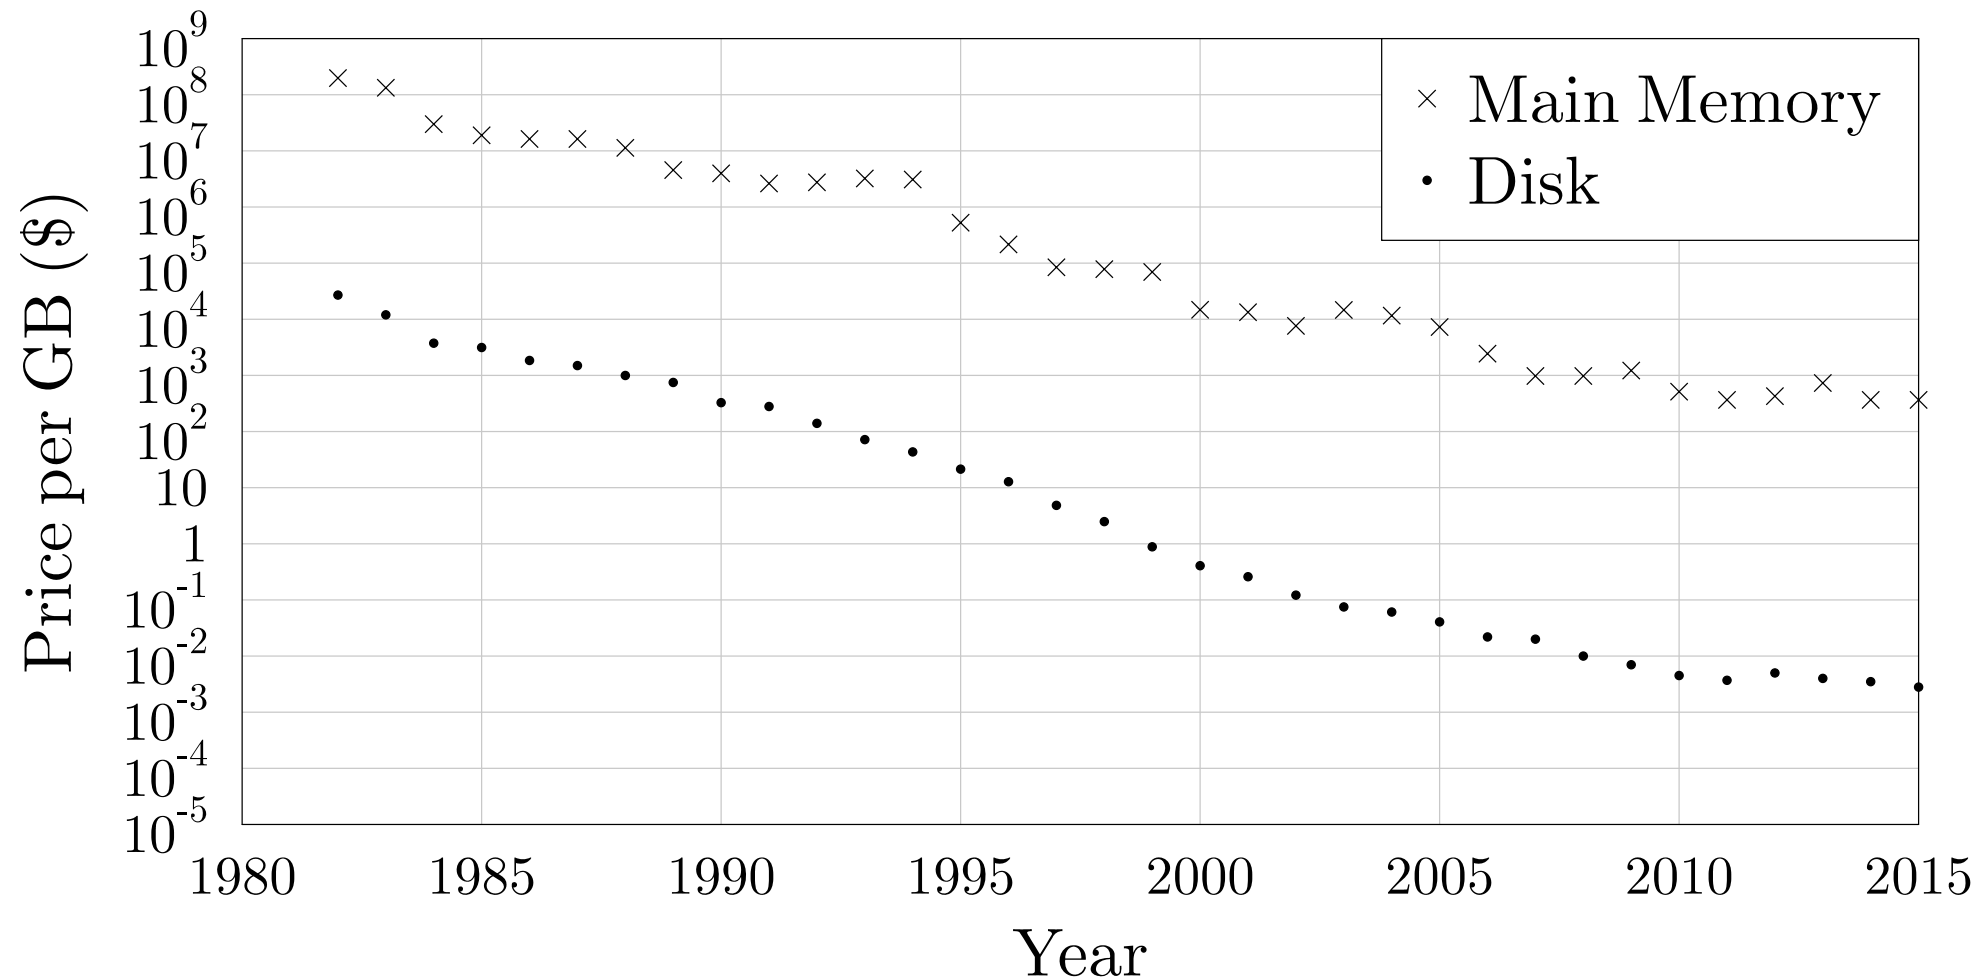
Results Catalogue

Better insert cost and **worse lookup cost** compared with B-trees

Can we improve the lookup cost?


	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Declining Main Memory Cost




Declining Main Memory Cost


Store a fence pointer for every block in main memory



Fence
pointers {



array {



Block 1	Block 2	Block 3	...
1	10	15	...
3	11	16	...
6	13	18	...

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(\log_2(N))$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N)^2)$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

Quick sanity check:

suppose
and

$$N = 2^{32}$$

$$B = 2^{10}$$

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree		
Tiered LSM-tree		

Results Catalogue – with fence pointers

Quick sanity check:


suppose
and

$$N = 2^{32}$$

$$B = 2^{10}$$

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(2^{31})$
Log	$O(2^{32})$	$O(2^{-10})$
B-tree	$O(4)$	$O(4)$
Basic LSM-tree	$O(32)$	$O(2^{-10} \cdot 32)$
Leveled LSM-tree		
Tiered LSM-tree		

Up until now we always create levels by merging **two** files!



Can we change that?

Leveled LSM-tree



Lookup cost




Update cost

Leveled LSM-tree

Lookup cost depends on number of levels
How to reduce it?

Increase size ratio T




Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



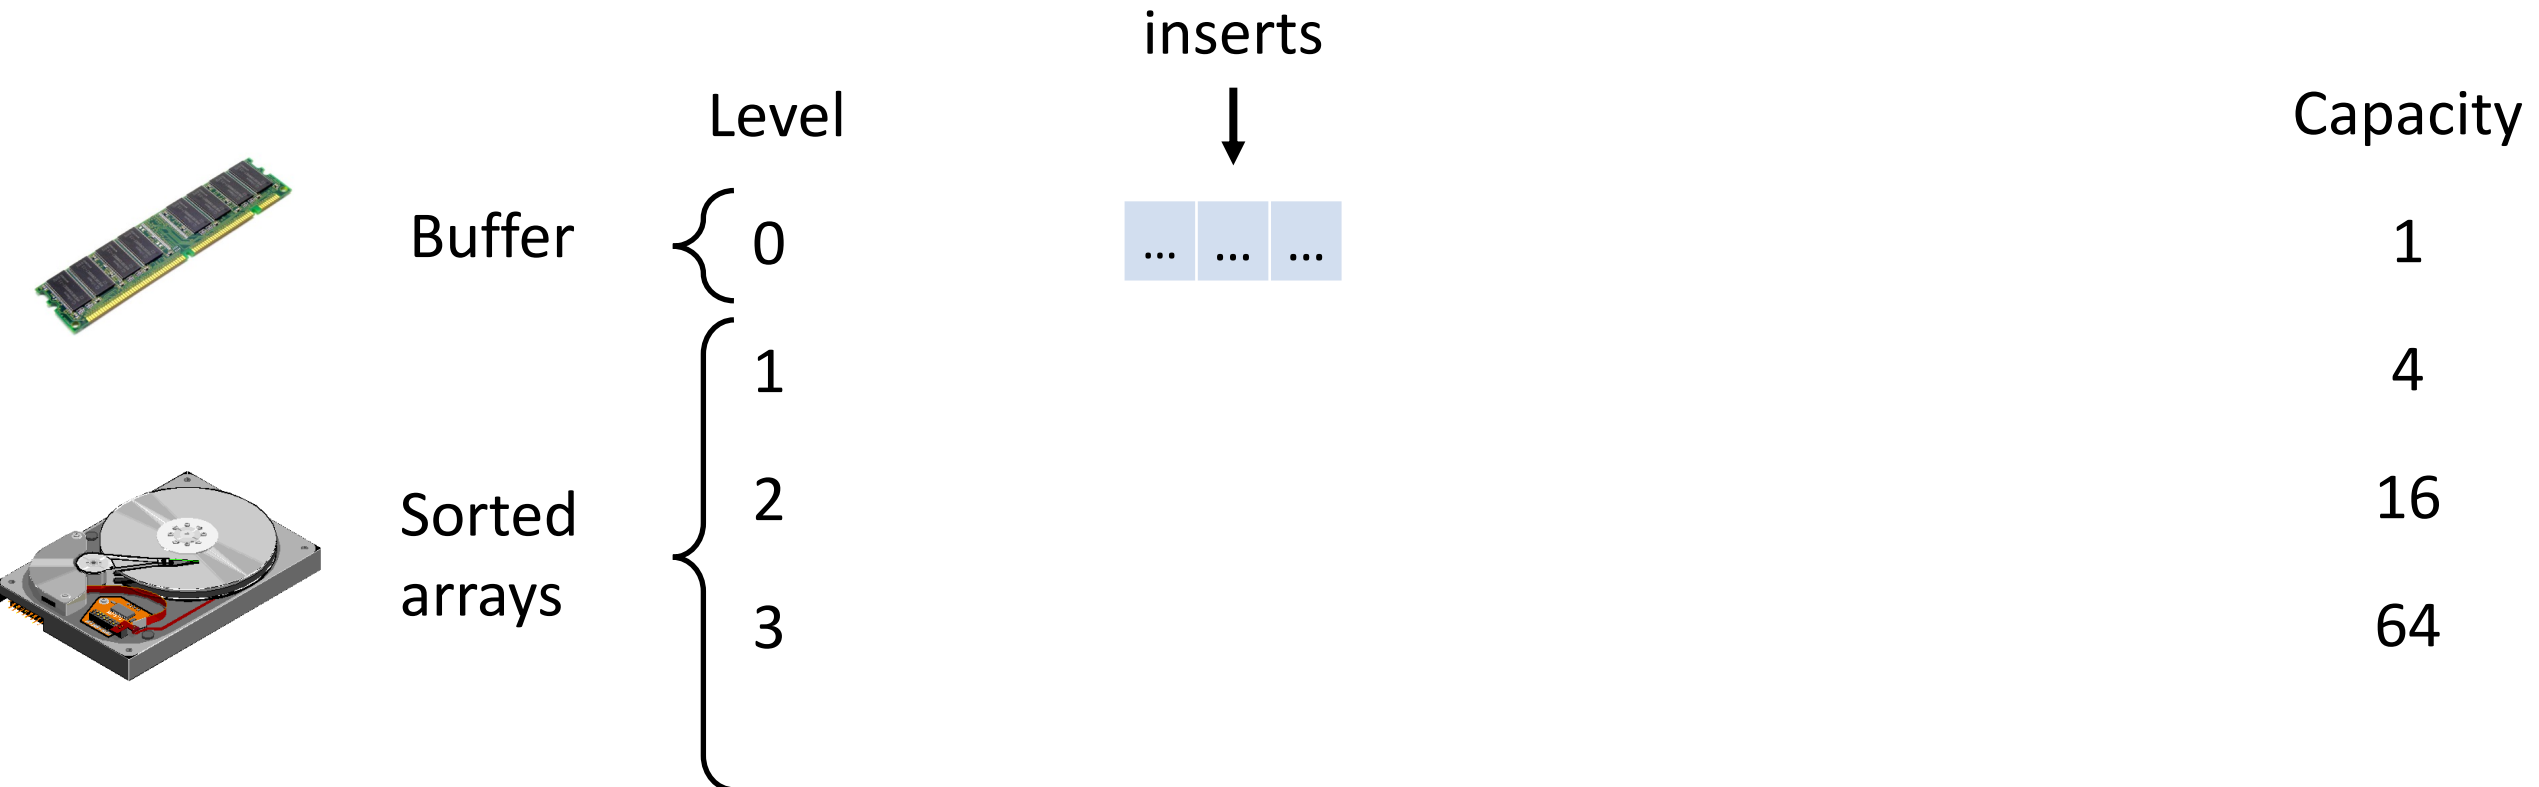
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T




Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



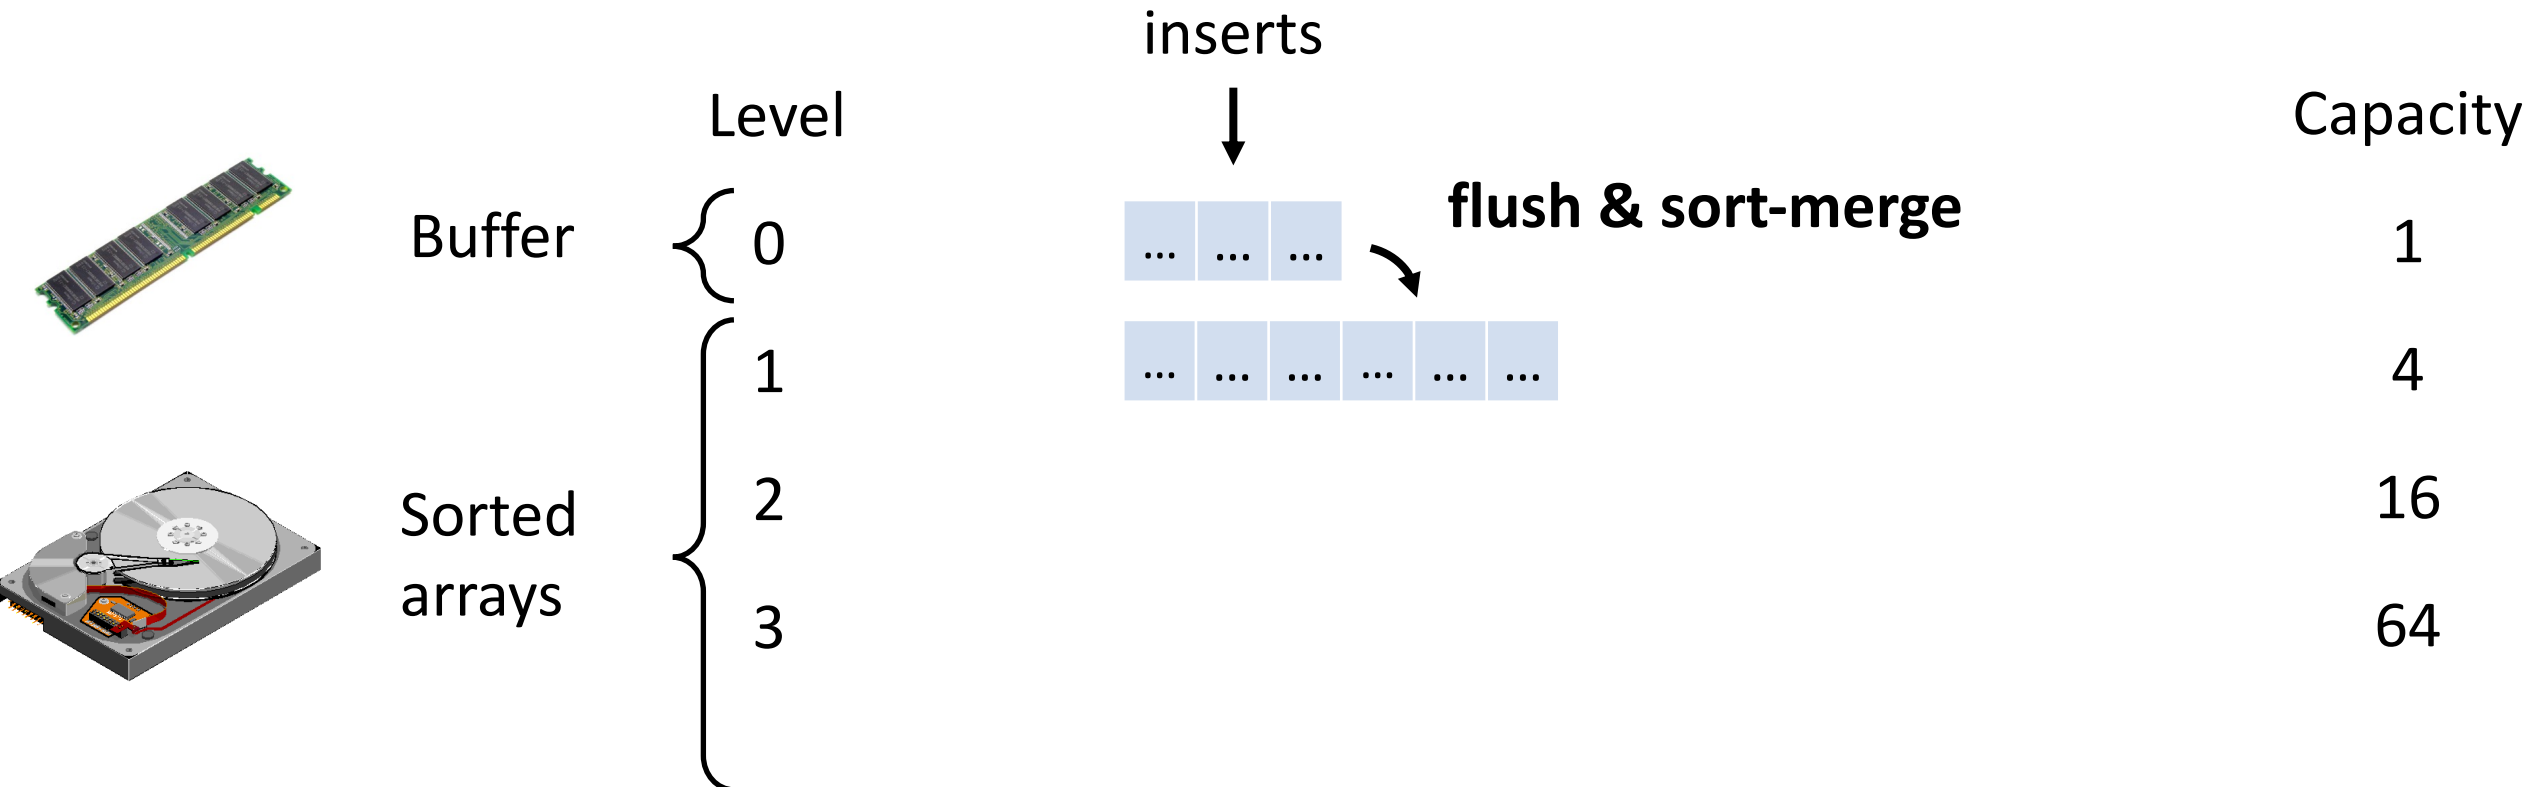
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



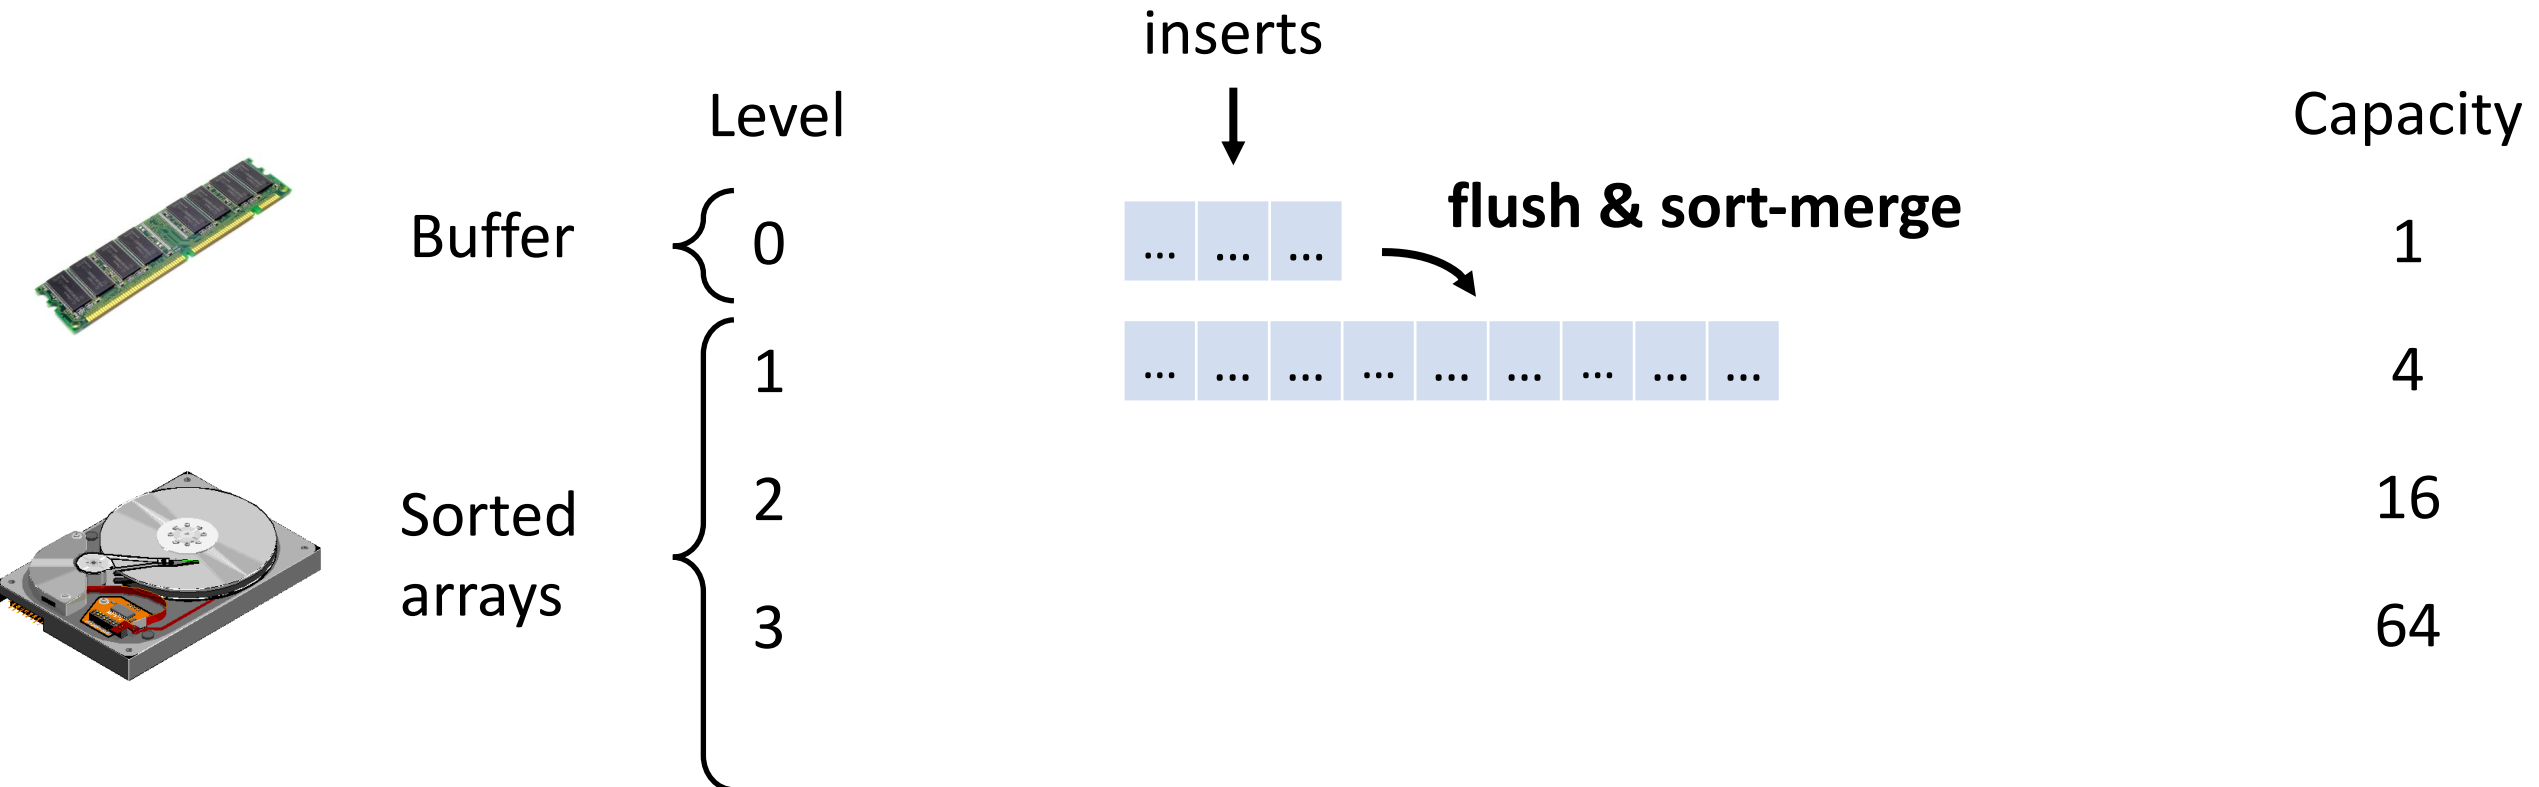
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



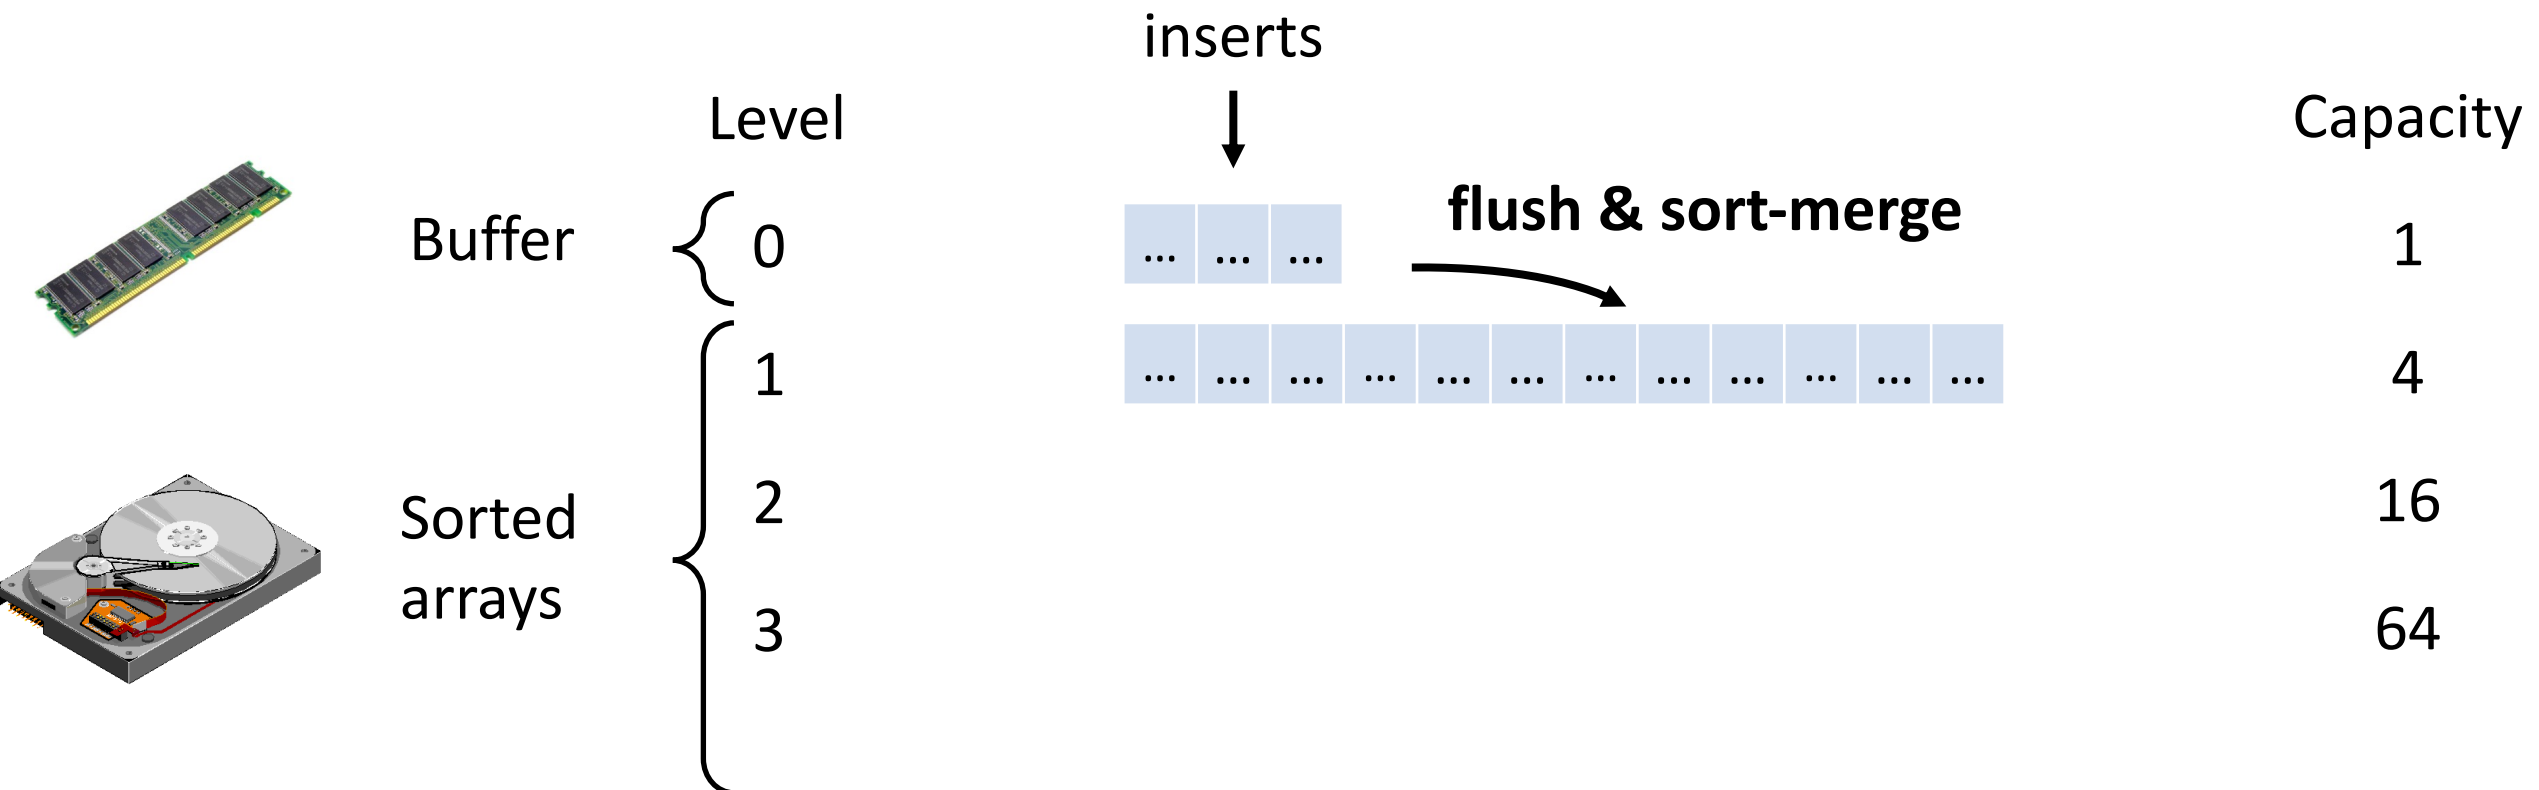
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



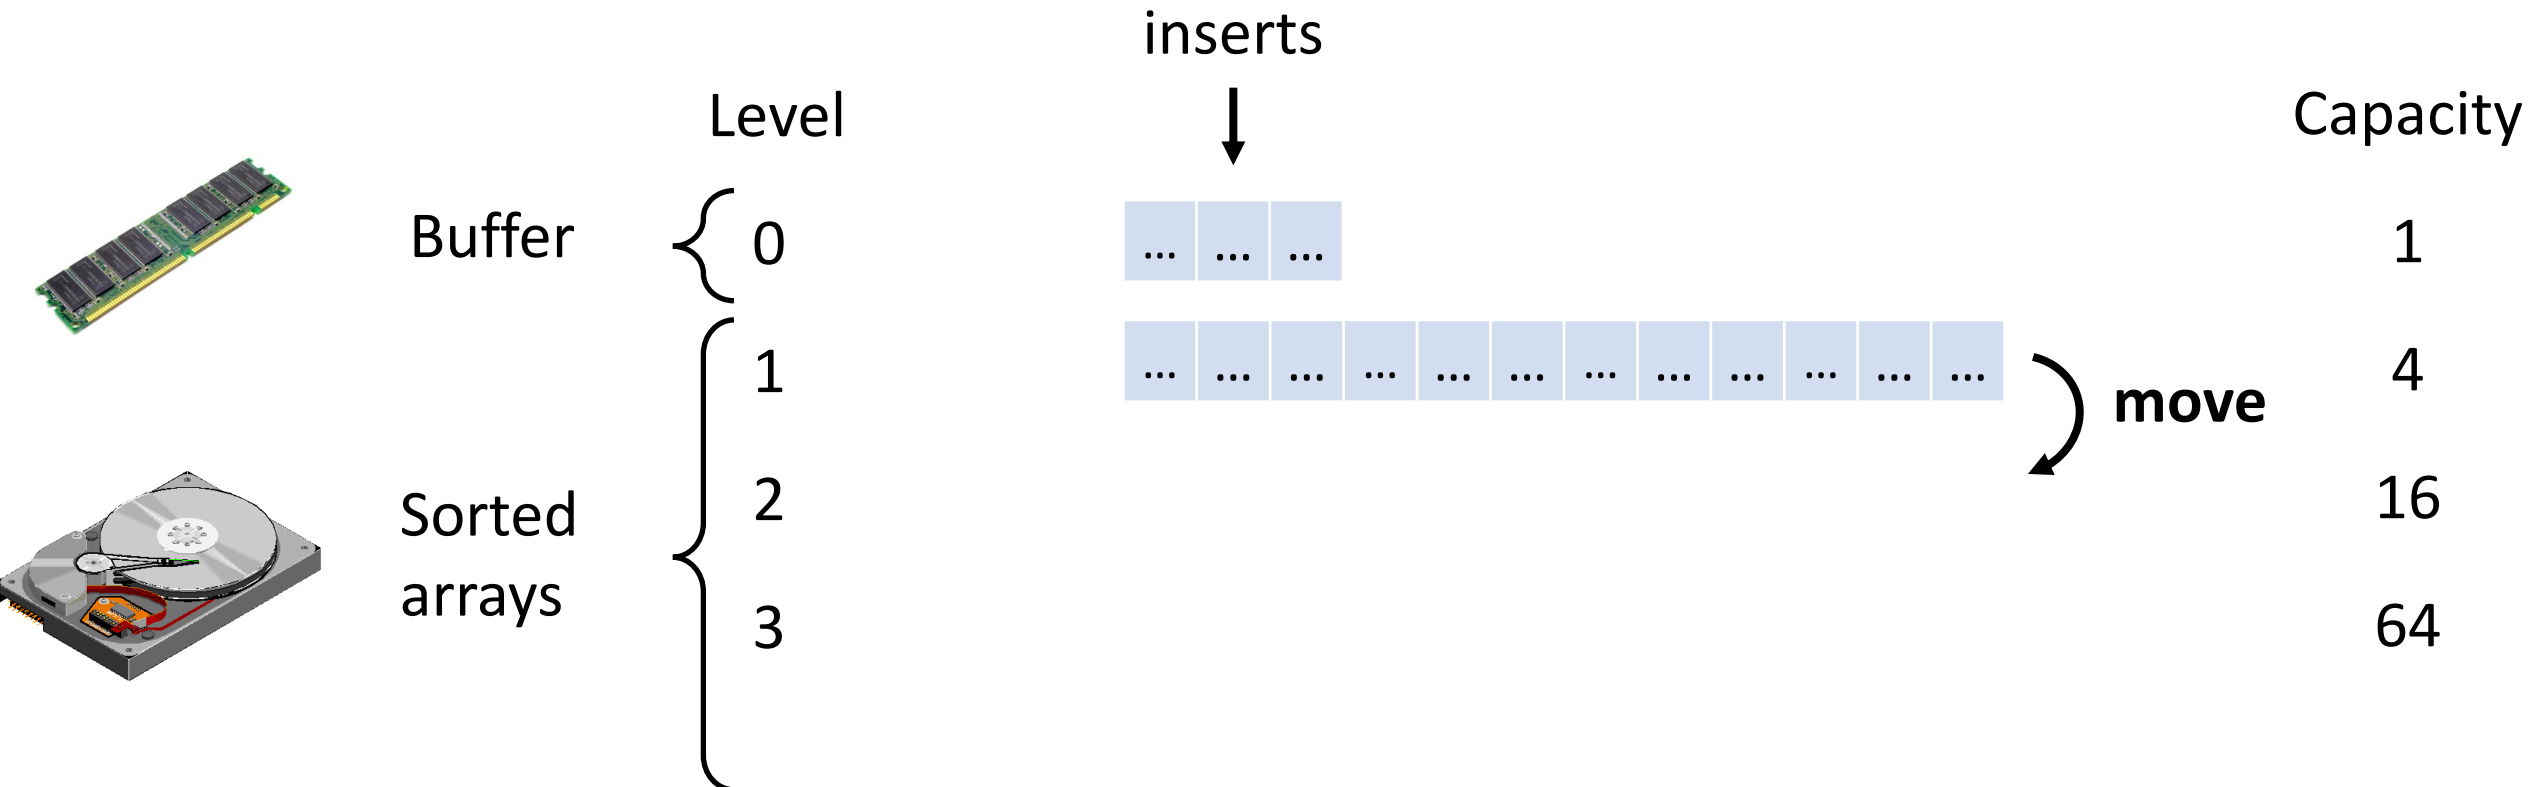
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



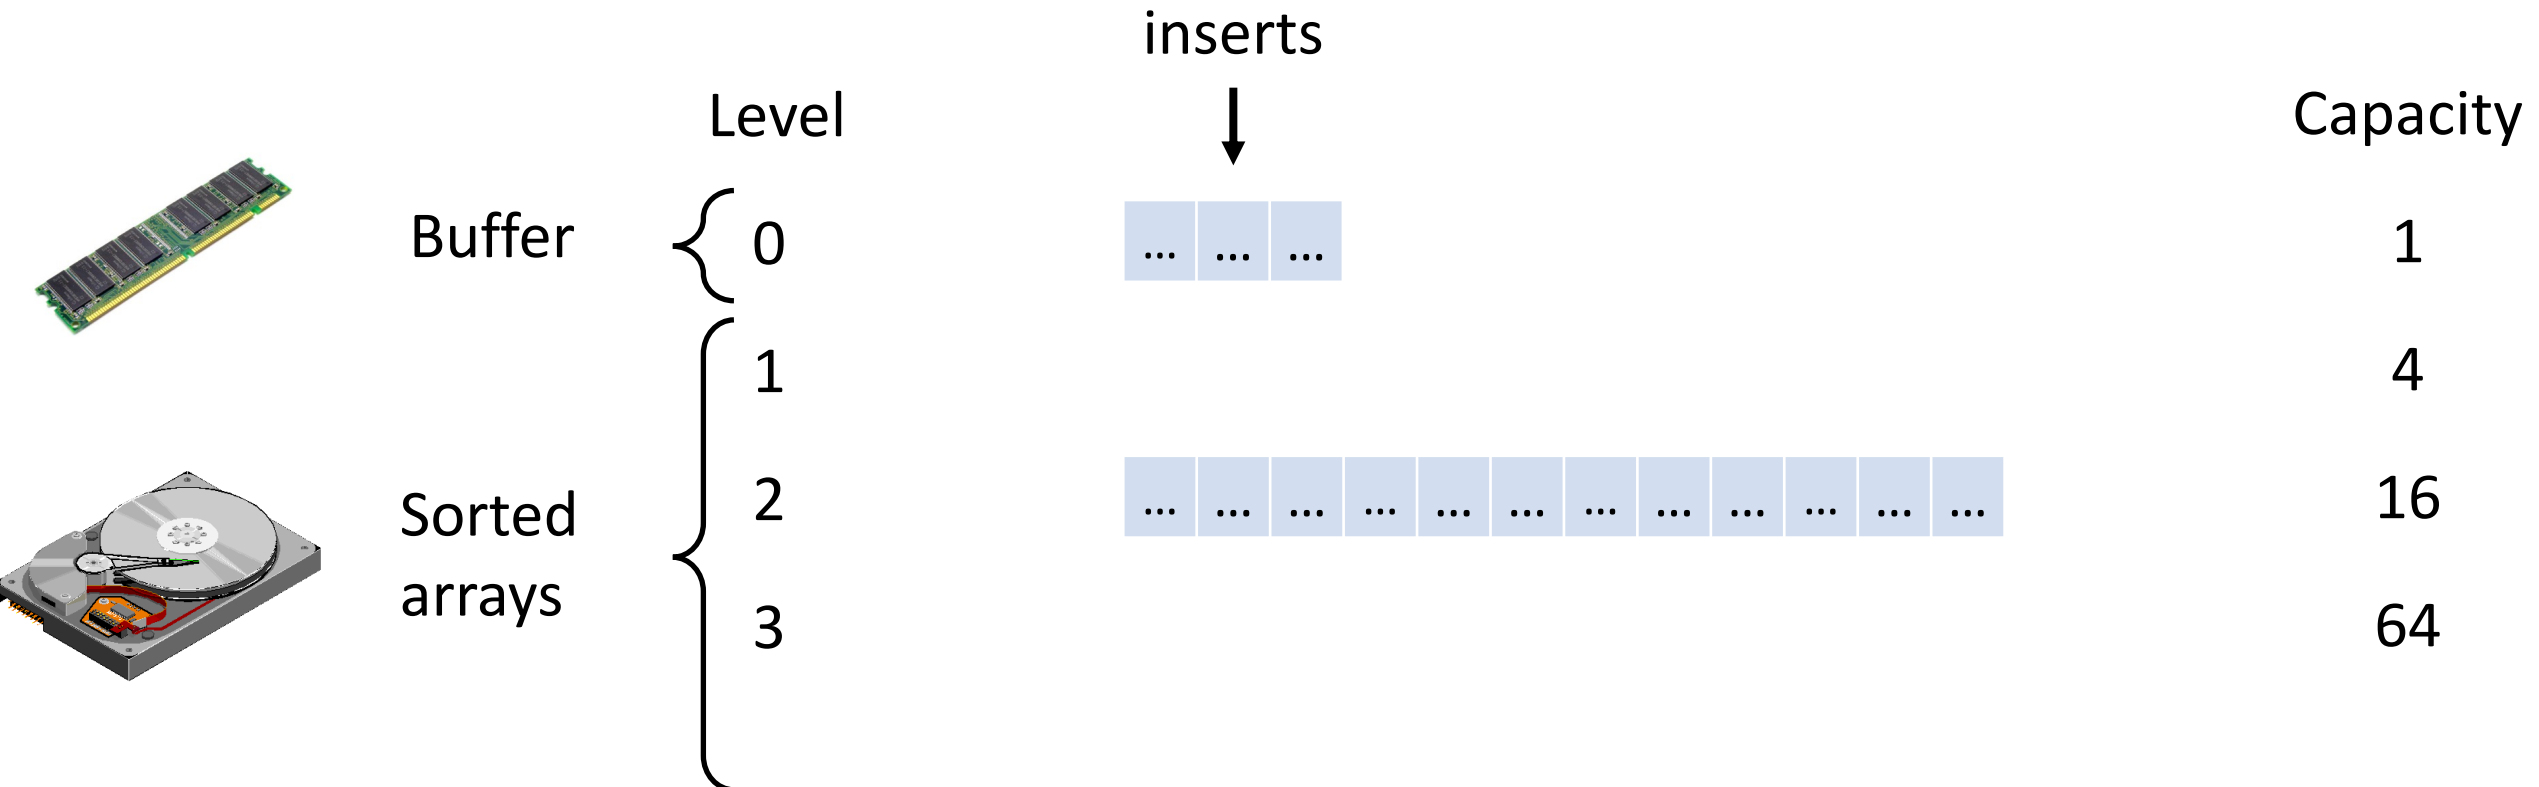
Leveled LSM-tree

Lookup cost depends on number of levels

How to reduce it?

E.g. size ratio of 4

Increase size ratio T



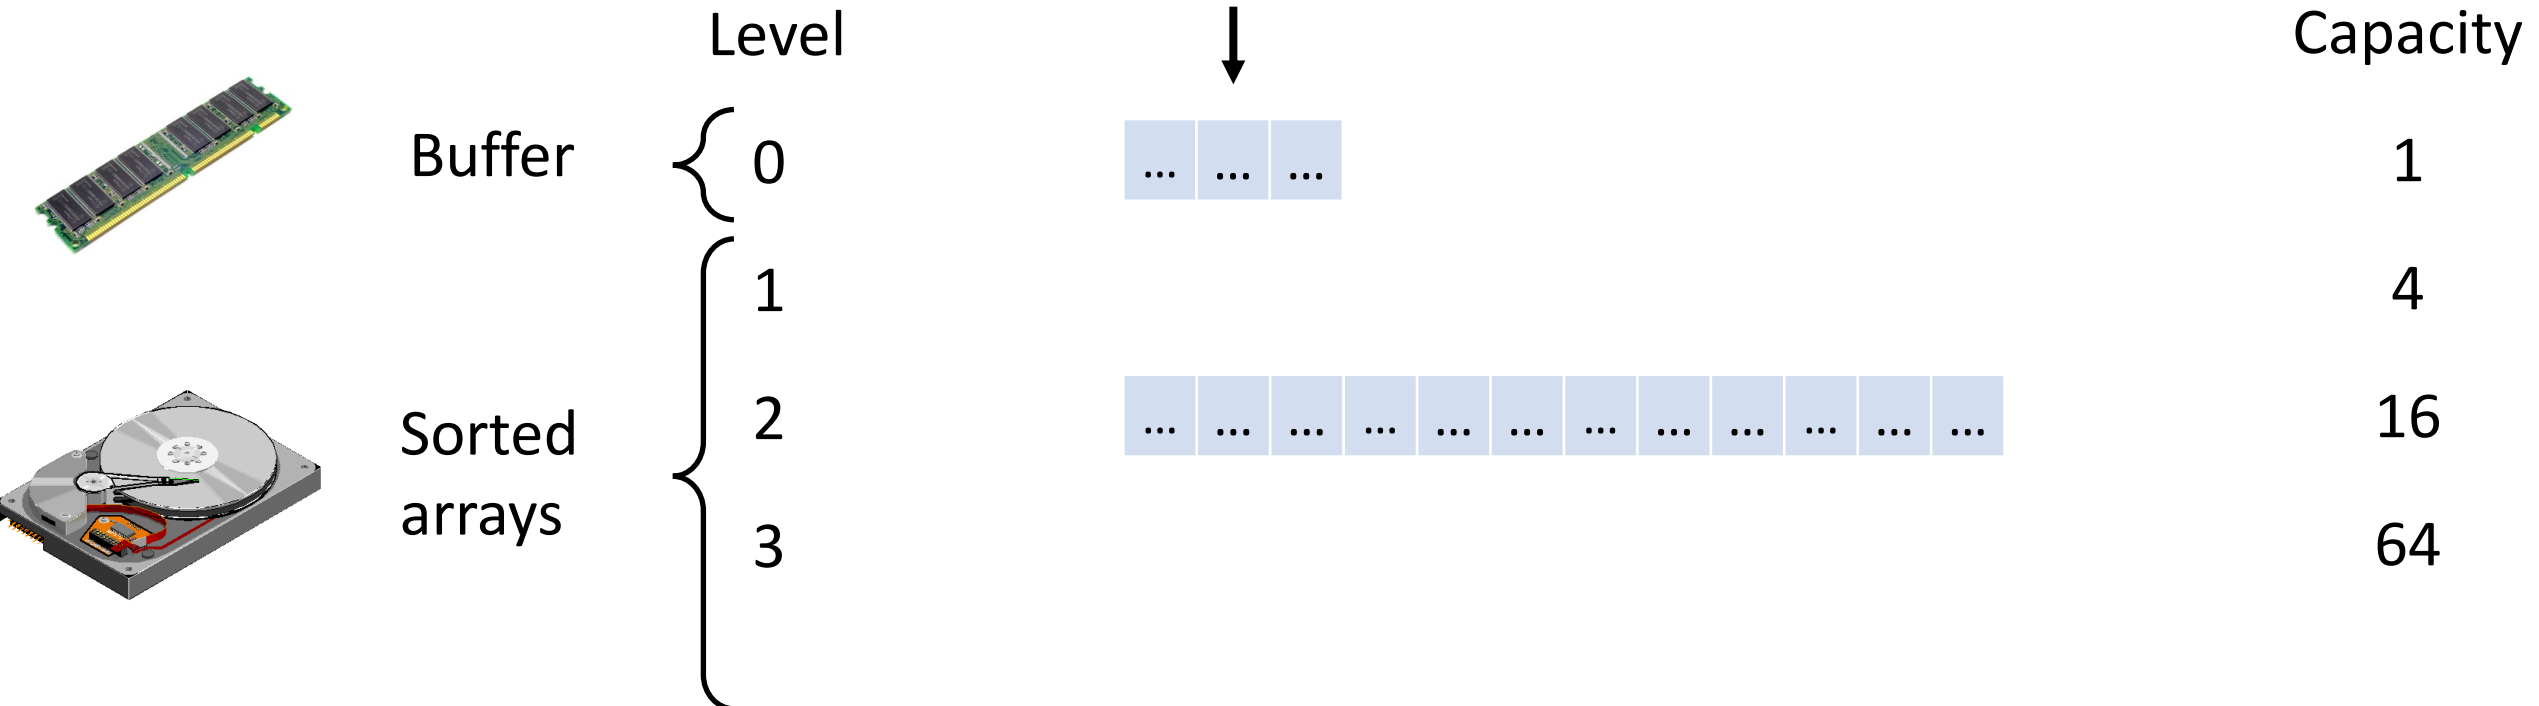
Leveled LSM-tree

Lookup cost?

$$O(\log_T(N))$$

Insertion cost?

$$O\left(\frac{T}{B} \cdot \log_T(N)\right)$$



Leveled LSM-tree

↓ Lookup cost?
 $O(\log_T(N))$

Insertion cost? ↑
 $O\left(\frac{T}{B} \cdot \log_T(N)\right)$

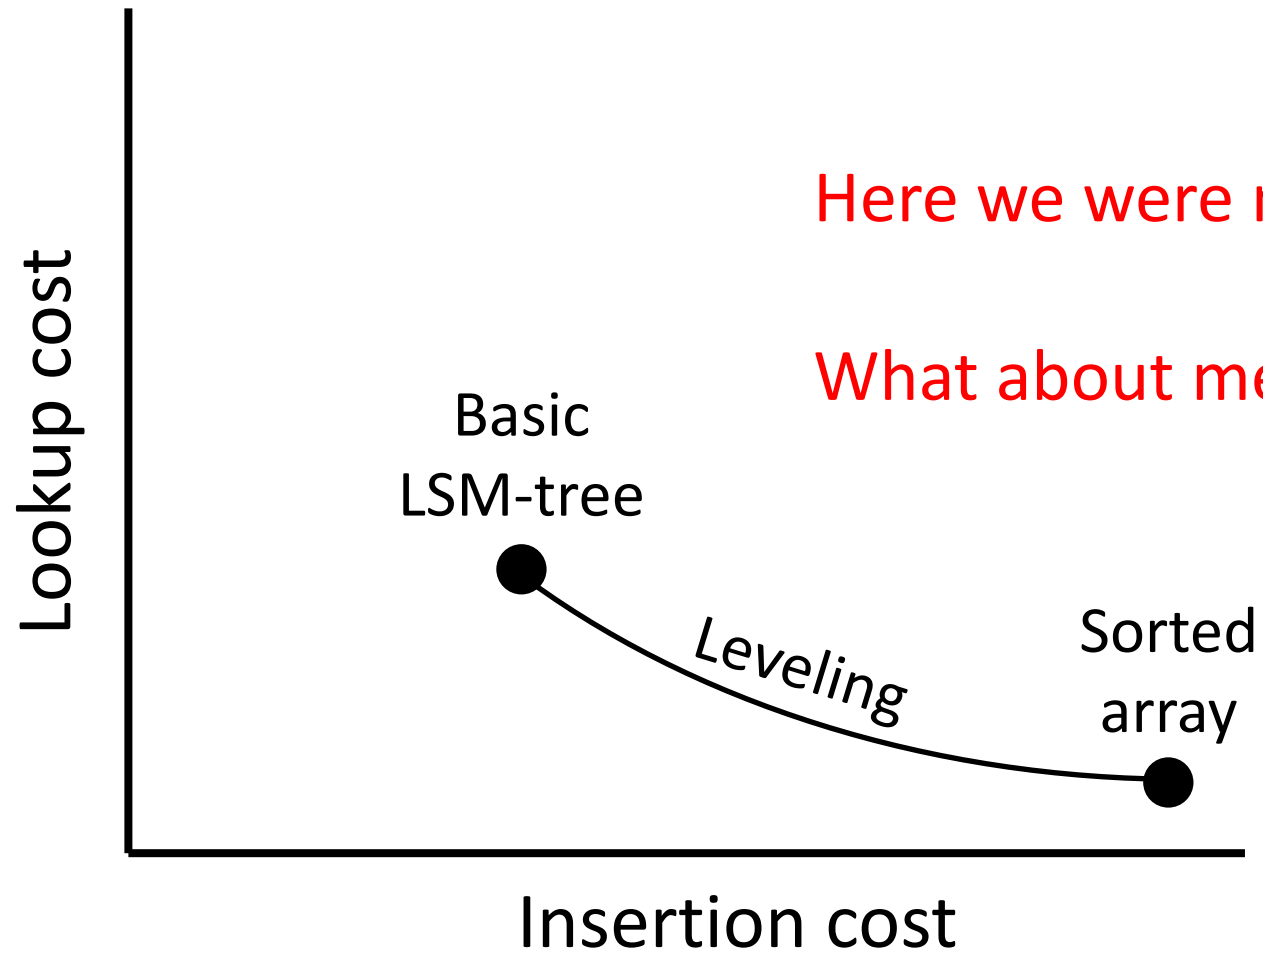
What happens as we increase the size ratio T ?

What happens when size ratio T is set to be N ?

Lookup cost becomes:
 $O(1)$

Insert cost becomes:
 $O(N/B)$

The LSM-tree becomes a sorted array!



Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree	$O(\log_T(N))$	$O(T/B \cdot \log_T(N))$
Tiered LSM-tree		

Tiered LSM-tree




Lookup cost



Insertion cost

Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.
Do not merge within a level.




Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4

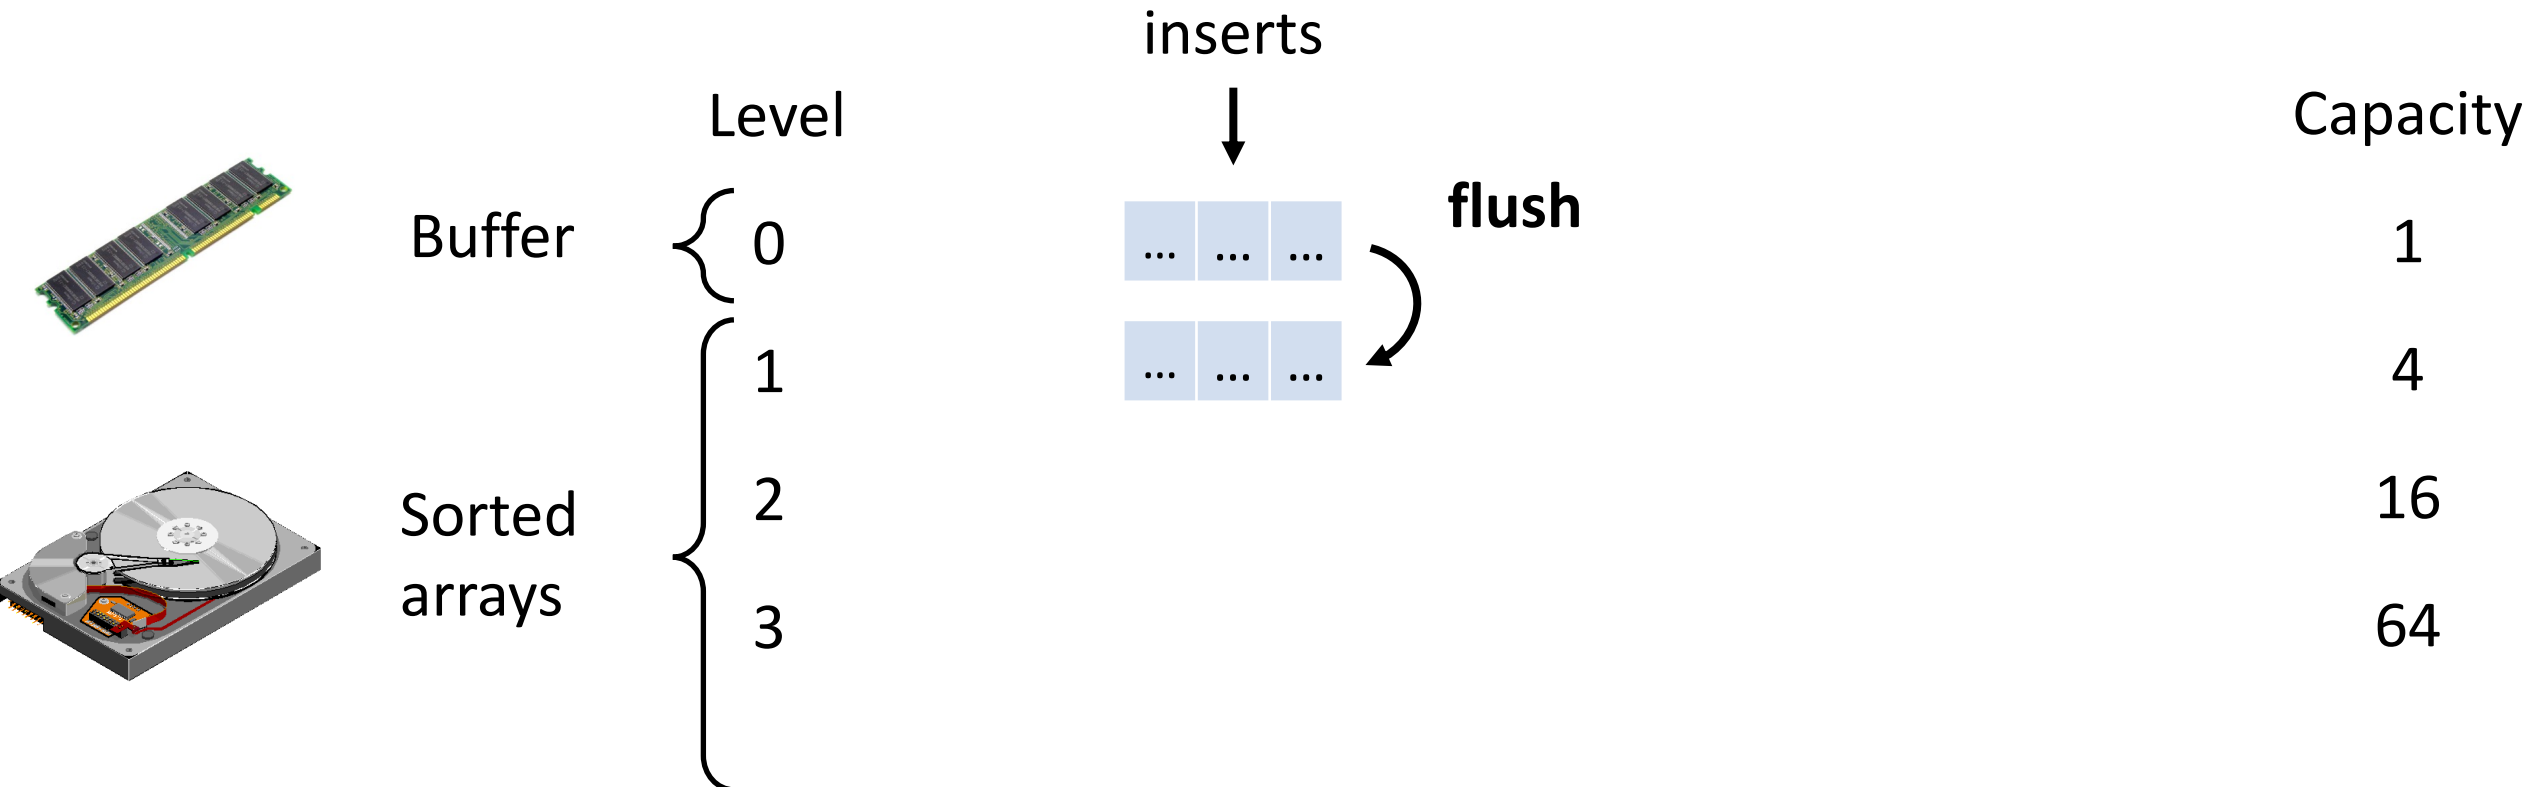


Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4

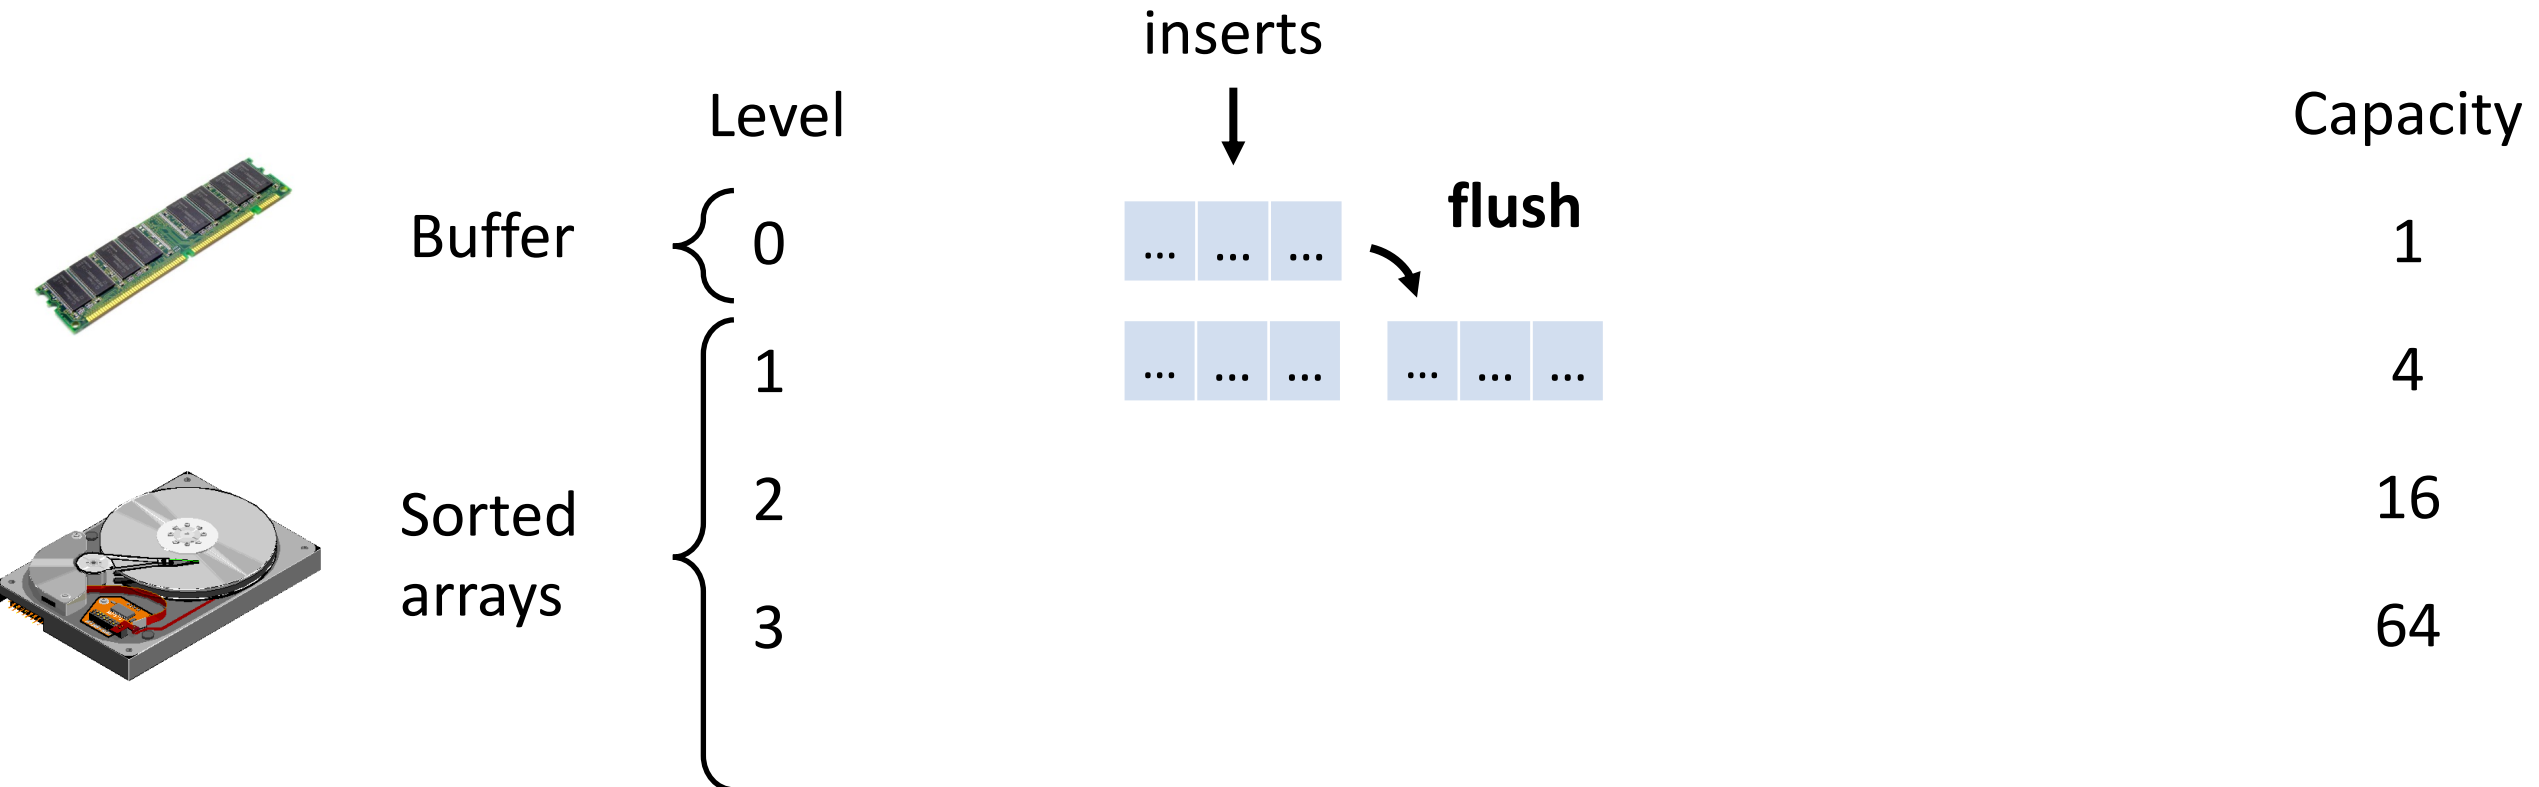


Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4

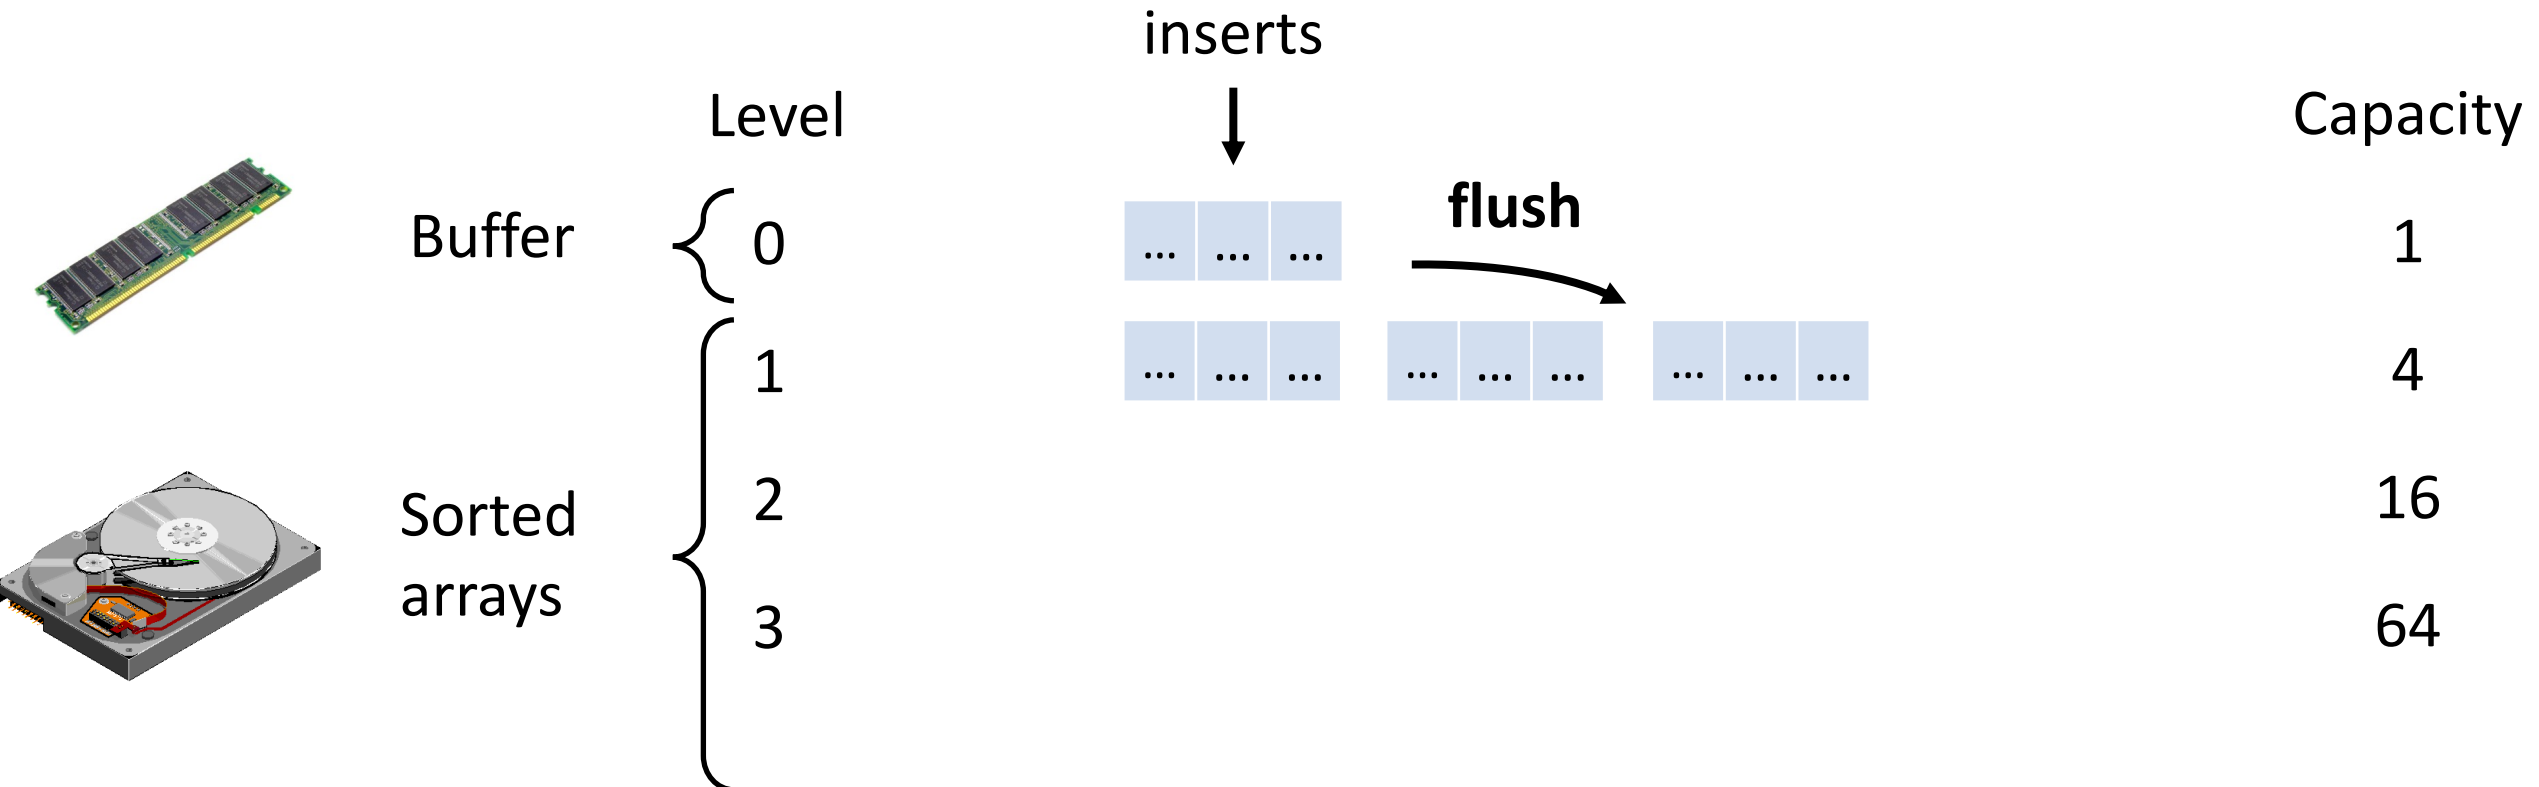


Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4




Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4

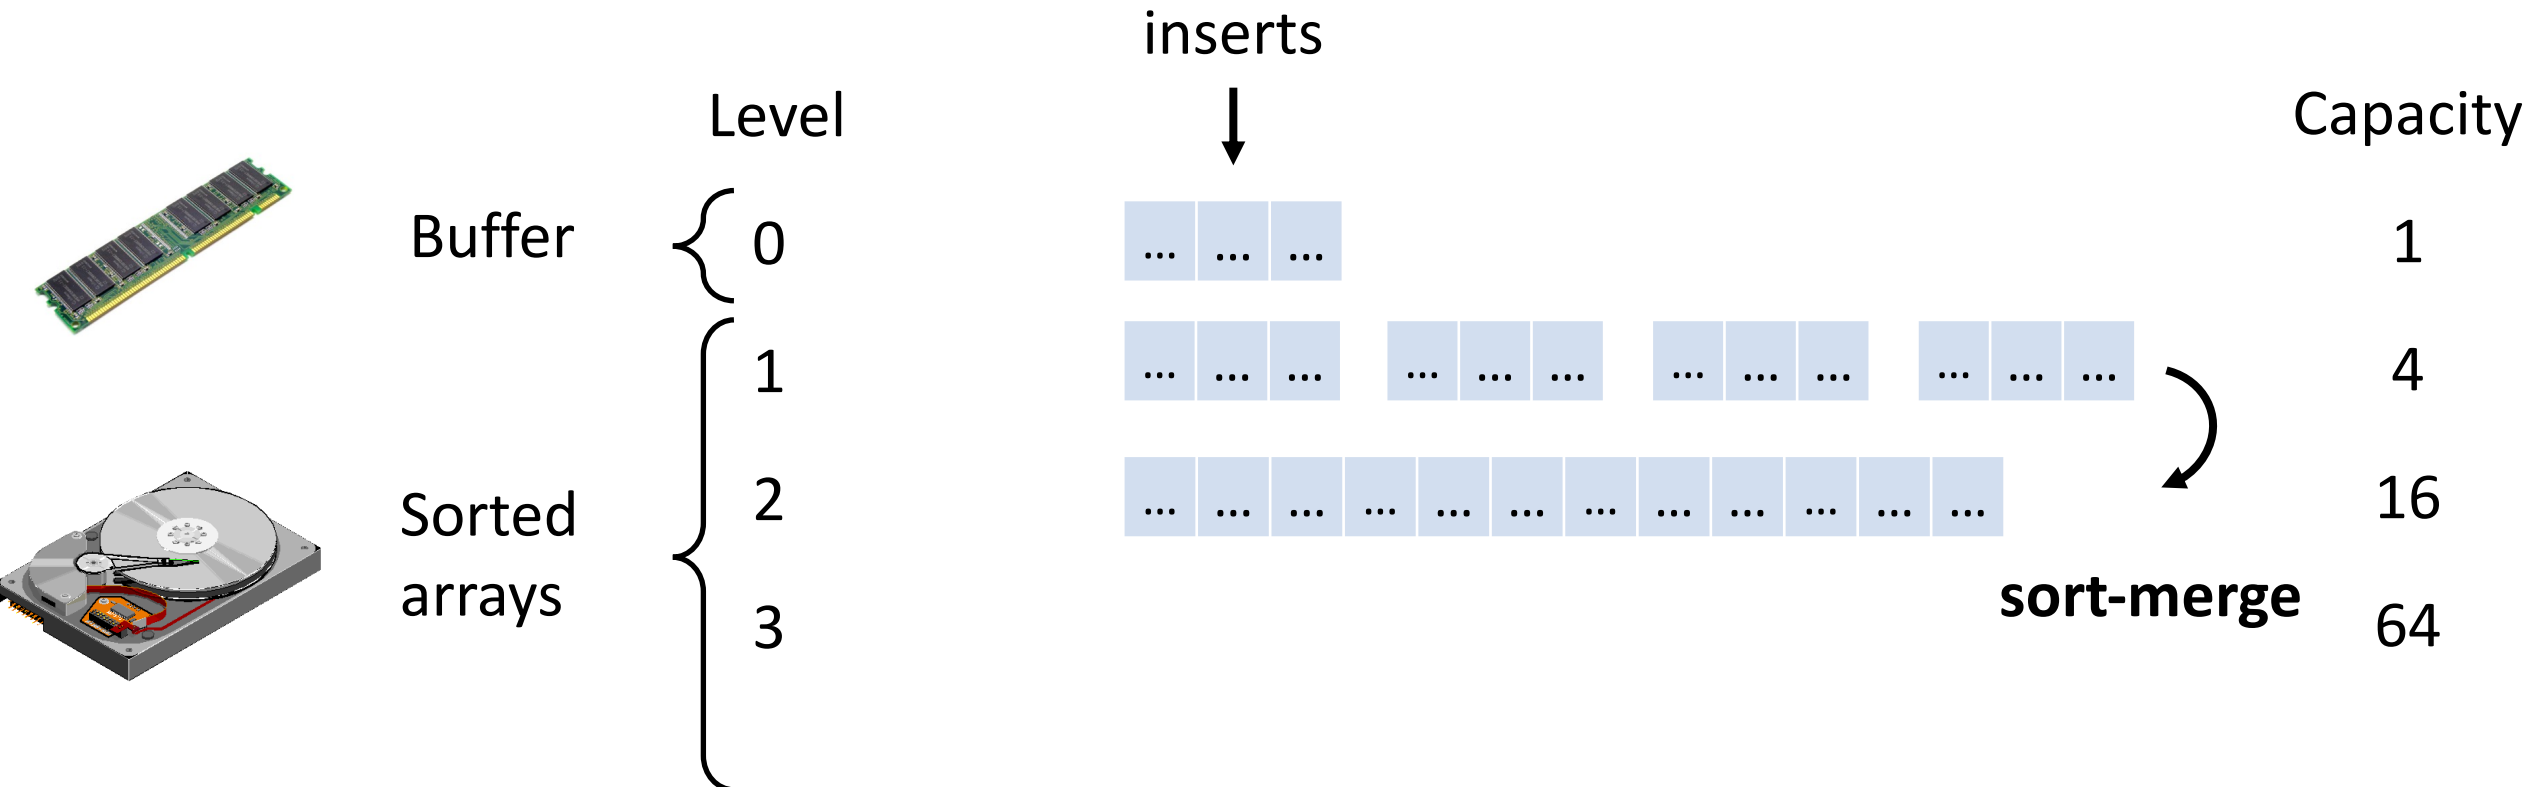


Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4




Tiered LSM-tree

Reduce the number of levels by increasing the size ratio.

Do not merge within a level.

E.g. size ratio of 4




Tiered LSM-tree

Lookup cost?

$$O(T \cdot \log_T(N))$$


Insertion cost?

$$O\left(\frac{1}{B} \cdot \log_T(N)\right)$$



Tiered LSM-tree

Lookup cost?
 $O(T \cdot \log_T(N))$

Insertion cost?
 $O\left(\frac{1}{B} \cdot \log_T(N)\right)$ 


What happens as we increase the size ratio T ?

What happens when size ratio T is set to be N ?

Lookup cost becomes:
 $O(N)$

Insert cost becomes:
 $O(1/B)$

The tiered LSM-tree becomes a log!



Results Catalogue – with fence pointers

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree	$O(\log_T(N))$	$O(T/B \cdot \log_T(N))$
Tiered LSM-tree	$O(T \cdot \log_T(N))$	$O(1/B \cdot \log_T(N))$

Results Catalogue – with fence pointers

Quick sanity check:

suppose

$$N = 2^{32}$$

and

$$B = 2^{10}$$

and

$$T = 2^2$$

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree	$O(\log_T(N))$	$O(T/B \cdot \log_T(N))$
Tiered LSM-tree	$O(T \cdot \log_T(N))$	$O(1/B \cdot \log_T(N))$

Results Catalogue – with fence pointers

Quick sanity check:

suppose

$$N = 2^{32}$$

and

$$B = 2^{10}$$

and

$$T = 2^2$$

	Lookup cost	Insertion cost
Sorted array	$2^0=1$	$2^{31}=2B$
Log	$2^{32}=4B$	$2^{-10}=0.001$
B-tree	$2^2=4$	$2^2=4$
Basic LSM-tree	$2^5=32$	$2^{-5}=0.031$
Leveled LSM-tree	$2^4=16$	$2^{-4}=0.063$
Tiered LSM-tree	$2^6=64$	$2^{-6}=0.016$

Results Catalogue – with fence pointers

Quick sanity check:

suppose

$$N = 2^{32}$$


and


$$B = 2^{10}$$

and

$$T = 10$$

	Lookup cost	Insertion cost
Sorted array	$2^0=1$	$2^{31}=2B$
Log	$2^{32}=4B$	$2^{-10}=0.001$
B-tree	$2^2=4$	$2^2=4$
Basic LSM-tree	$2^5=32$	$2^{-5}=0.031$
Leveled LSM-tree	$\log_{10}(2^{32})=9.6$	$10 \cdot 2^{-10} \cdot \log_{10}(2^{32}) = 0.09$
Tiered LSM-tree	$10 \cdot \log_{10}(2^{32})=96$	$2^{-10} \cdot \log_{10}(2^{32}) = 0.009$





How to avoid them?

An **oracle** that helps us to skip them!

Bloom filters

Answer **set-membership** queries

Small size, typically stored in **memory**

May return **false positives**


Bloom filters

k hash functions

$h_1(\blacksquare)$


$h_2(\blacksquare)$

$h_3(\blacksquare)$




an array of m bits


Bloom filters – insert v_1




Bloom filters – insert v_2



Bloom filters – query v_3



Bloom filters – query v_4




false positive rate: $f = e^{-\frac{m}{n} \cdot (\ln(2))^2}$

sanity check: for $\frac{m}{n} = 10$, $f = 0.00819$


after inserting n elements

→ we have m/n **bits per key**

Augmenting the LSM design with Bloom filters



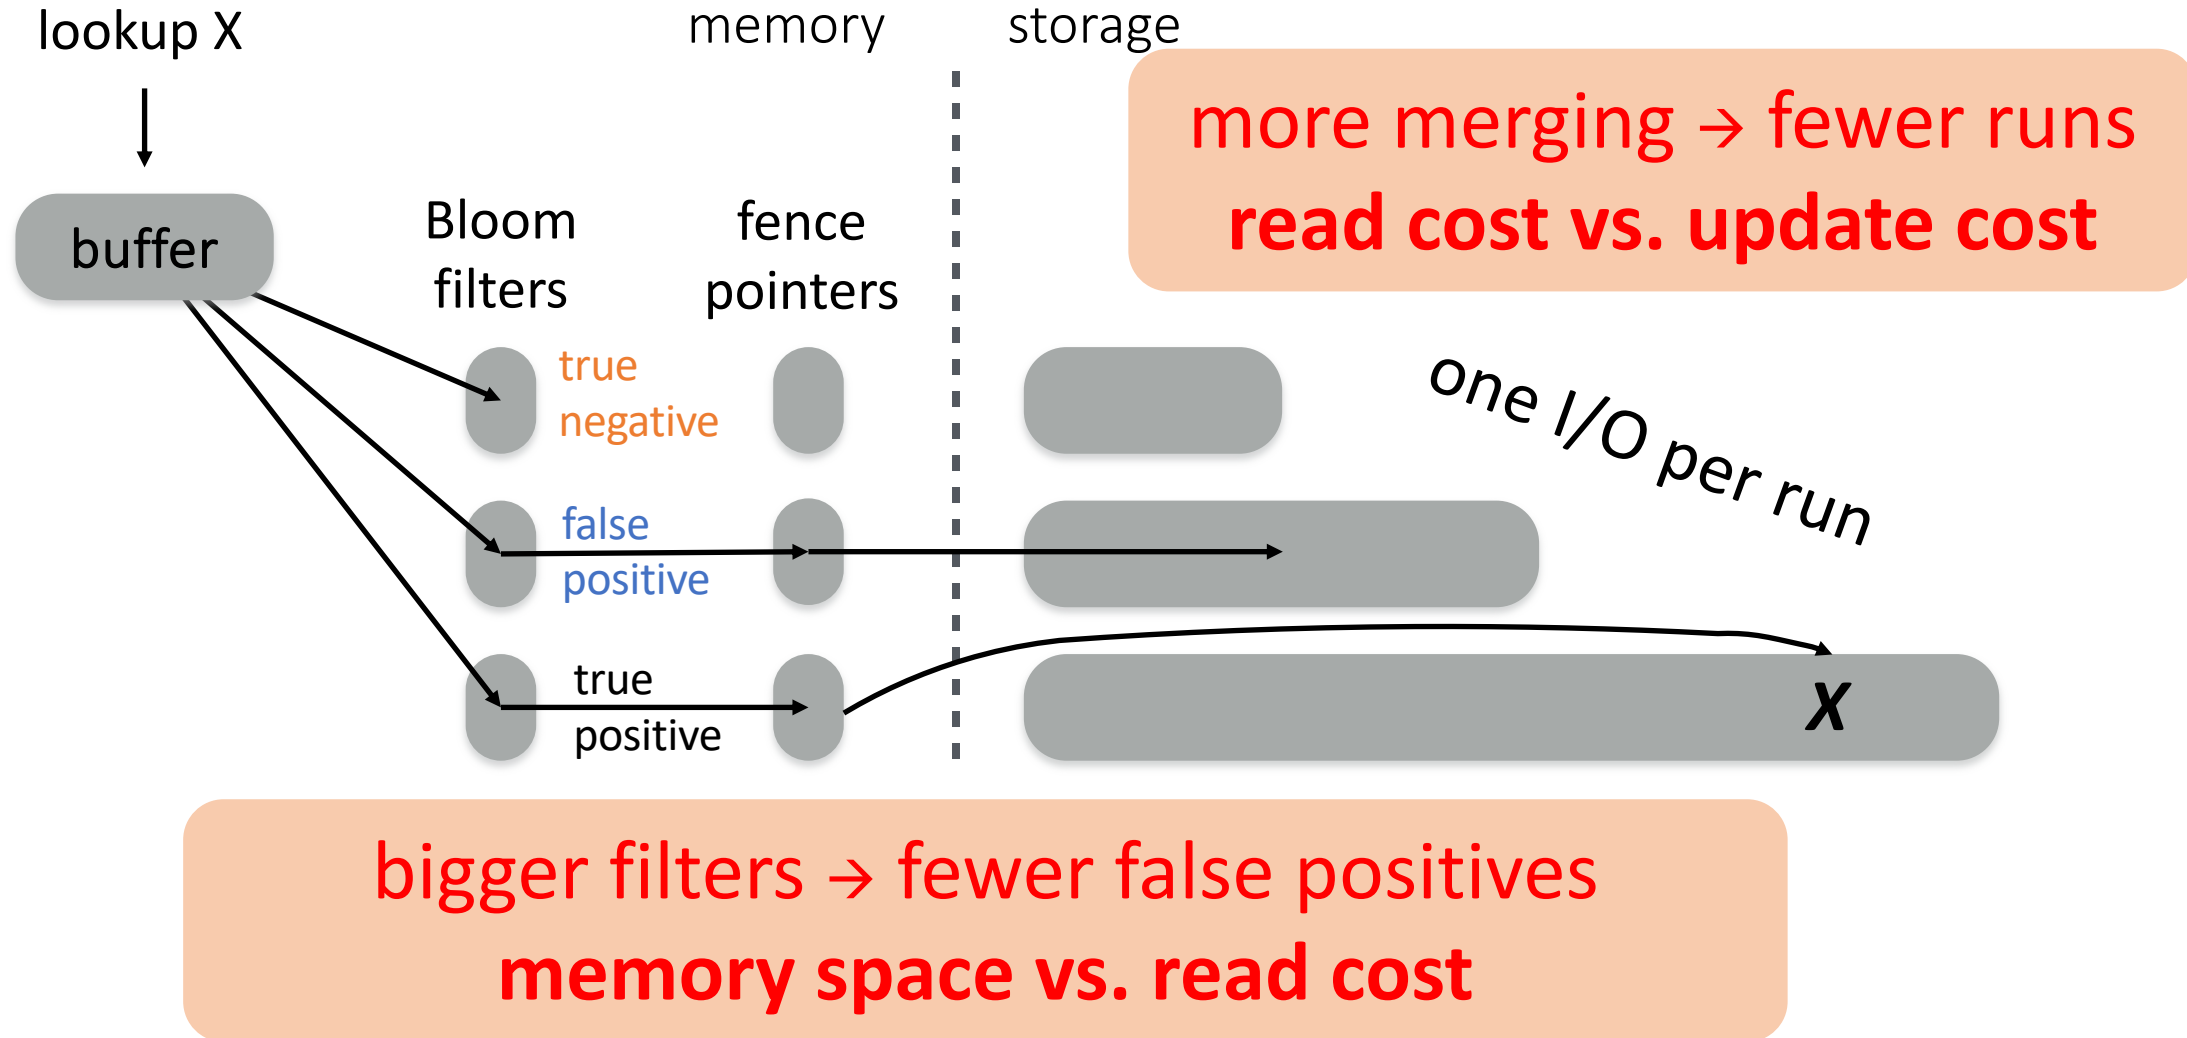
Augmenting the LSM design with Bloom filters



Empty Queries: only FPs

Non-Empty Queries: FPs and one I/O

performance & cost trade-offs



Results Catalogue – with fence pointers & BFs

Empty Queries

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree	$O(f \cdot \log_T(N))$	$O(T/B \cdot \log_T(N))$
Tiered LSM-tree	$O(f \cdot T \cdot \log_T(N))$	$O(1/B \cdot \log_T(N))$

Results Catalogue – with fence pointers & BFs

Quick sanity check:

suppose

$$N = 2^{32}$$

and

$$B = 2^{10}$$

and

$$T = 10 \text{ and } m/n = 10$$

Empty Queries

	Lookup cost	Insertion cost
Sorted array	$2^0=1$	$2^{31}=2B$
Log	$2^{32}=4B$	$2^{-10}=0.001$
B-tree	$2^2=4$	$2^2=4$
Basic LSM-tree	$2^5=32$	$2^{-5}=0.031$
Leveled LSM-tree	$f \cdot \log_{10}(2^{32})=0.079$	$10 \cdot 2^{-10} \cdot \log_{10}(2^{32}) = 0.09$
Tiered LSM-tree	$f \cdot 10 \cdot \log_{10}(2^{32})=0.79$	$2^{-10} \cdot \log_{10}(2^{32}) = 0.009$

Results Catalogue – with fence pointers & BFs

Non-Empty Queries

	Lookup cost	Insertion cost
Sorted array	$O(1)$	$O(N/2)$
Log	$O(N)$	$O(1/B)$
B-tree	$O(\log_B(N))$	$O(\log_B(N))$
Basic LSM-tree	$O(\log_2(N))$	$O(1/B \cdot \log_2(N))$
Leveled LSM-tree	$O(1 + f \cdot \log_T(N))$	$O(T/B \cdot \log_T(N))$
Tiered LSM-tree	$O(1 + f \cdot T \cdot \log_T(N))$	$O(1/B \cdot \log_T(N))$

Results Catalogue – with fence pointers & BFs

Quick sanity check:

suppose

$$N = 2^{32}$$

and

$$B = 2^{10}$$


and

$$T = 10 \text{ and } m/n = 10$$

Non-Empty Queries

	Lookup cost	Insertion cost
Sorted array	$2^0=1$	$2^{31}=2B$
Log	$2^{32}=4B$	$2^{-10}=0.001$
B-tree	$2^2=4$	$2^2=4$
Basic LSM-tree	$2^5=32$	$2^{-5}=0.031$
Leveled LSM-tree	$1 + f \cdot \log_{10}(2^{32})=1.079$	$10 \cdot 2^{-10} \cdot \log_{10}(2^{32}) = 0.09$
Tiered LSM-tree	$1 + f \cdot 10 \cdot \log_{10}(2^{32})=1.79$	$2^{-10} \cdot \log_{10}(2^{32}) = 0.009$

Bloom Filters



Conclusions

Write-optimized

Highly tunable

Backbone of many modern systems

Trade-off between lookup and insert cost (tiering/leveling, size ratio)

Trade main memory for lookup cost (fence pointers, Bloom filters)

Thank you!